

# Partner Execution Guide: Innovate with AI Apps and Agents

Revision: October 2025

**NOTE: This Guide is for Microsoft  
partner audiences only.**



# Topics Included

- Routes to Success
  - Routes to Success with customers
  - Key Use-cases
  - Win Formula (*Must review!*)
  - Key Actions (*Must review!*)
- Drive Apps that Matter and establish AI Foundry as the platform of choice
  - Sales Guidance (*Must review!*)
  - Product highlights
- Connect with developers with GitHub and driving them to Azure
  - Sales Guidance (*Must review!*)
  - Product highlights
- Closing Strong
  - Azure Accelerate
  - Azure Essentials

# What is an AI frontier firm



## The Frontier Firm Defined

A next-generation company that fuses human judgment with AI agents to scale faster, work smarter, and unlock new value - powered by on-demand intelligence and human + agent hybrid teams.



## Core Characteristics

- **AI-operated, human-led:** Agents execute tasks; humans provide oversight.
- **Work Chart vs Org Chart:** Teams form around outcomes, not departments.
- **Agent boss mindset:** Every employee manages and collaborates with AI agents.



## Why This Matters Now

- **82%** of leaders say 2025 is a pivotal year to rethink strategy and operations.\*
- **80%** expect agents to be integrated into their AI strategy within 12–18 months.\*
- Frontier Firms outperform peers in productivity, innovation, and speed to value

"As AI democratizes access to expertise and intelligence, we'll see the rise of Intelligence Resources departments -much like how HR and IT evolved into core functions. These new divisions will be essential for managing the interplay between humans and AI agents, emerging as a critical source of competitive advantage in the AI-enabled enterprise."

-Karim R. Lakhani, Chair of the Digital, Data, and Design Institute at Harvard,

[Microsoft Work Trends Index Report](#)



# Why innovate with Azure AI apps and agents

**Empower developers to easily build AI apps and agents**

Enable developers **with integrated AI-native tools**

Reduce developer ramp-up and unlock AI-driven innovation across teams

**Design, customize, and manage AI apps and agents at scale**

Orchestrate models, agents, and AI apps through a **unified platform**

Get end-to-end visibility and centralized resource management across departments

**Scale AI apps and agents seamlessly and efficiently**

Streamline **deployment and scaling of AI services** with enterprise-grade databases and app services

Optimize resource utilization, run AI apps with high performance, and align AI scaling with ROI goals

**Secure, govern, and ensure the safety of AI**

Mitigate risk and ensure **security and governance** with powerful tools

Empower responsible AI practices with built-in security, governance, and observability

# Journey to the Frontier Firm

## Phase 1

Human with assistant



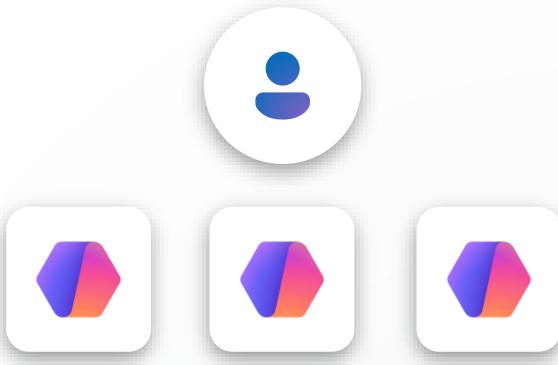
+



Every employee has an AI assistant that helps them work better and faster

## Phase 2

Human-led agents



Agents join teams as “digital colleagues,” taking on specific tasks at human direction

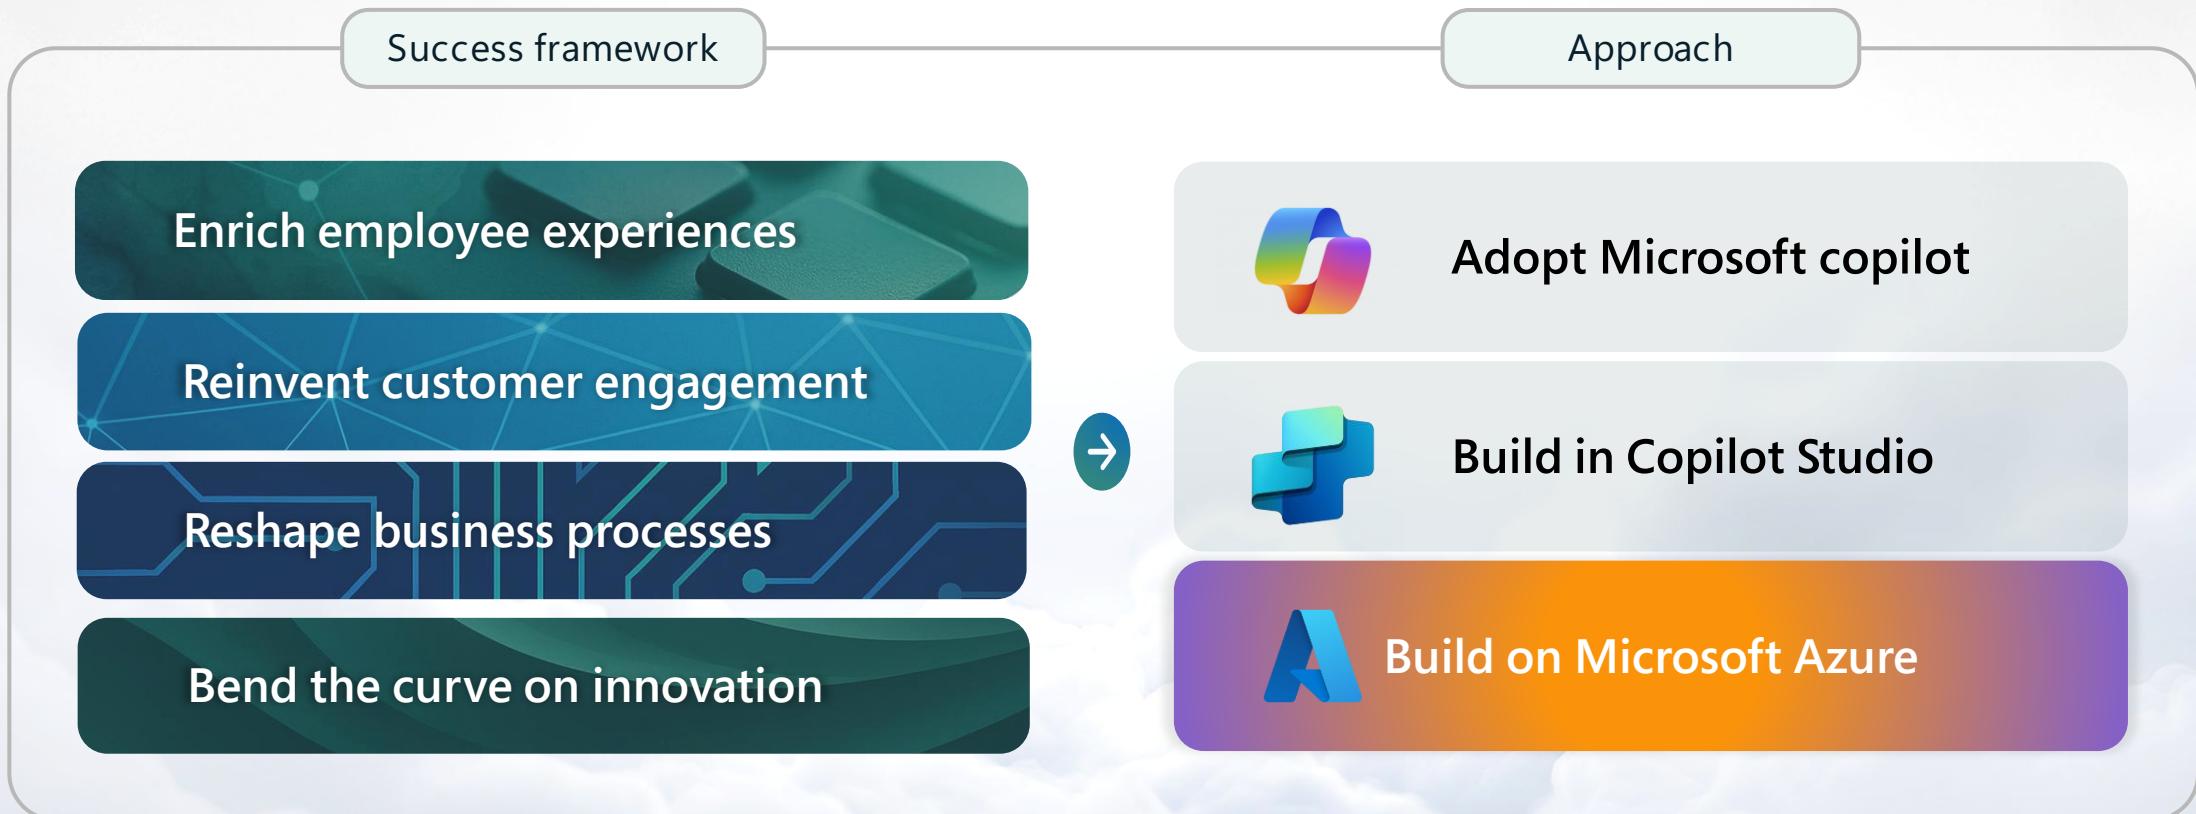
## Phase 3

Human-led, agent-operated



Humans set direction and agents run entire business processes and workflows, checking in as needed

# Becoming Frontier



Focus of this playbook

# Routes to Success



# FY26 Routes to Success with Customers

Drive Apps that Matter  
and establish  
AI Foundry as the  
platform of choice

Connect with  
developers with GitHub  
and drive them to  
Azure

# Key Use-Cases: Drive AI impact based on your industry

## Banking

- Smart document processing and insights
- AI-powered contact centers and support
- Proactive fraud detection and compliance
- Streamlined banking operations

## Capital Markets

- AI summaries of reports and earnings calls
- Market sentiment analysis
- Enhanced accessibility (such as audio descriptions)

## Insurance

- AI-assisted policy and document management
- Smarter agent support and sales enablement
- Compliance reviews and complex claims automation

## Retail

- Personalized shopping assistants
- AI-driven market insights
- Intelligent contact center support
- Automated retail operations

## Consumer Goods

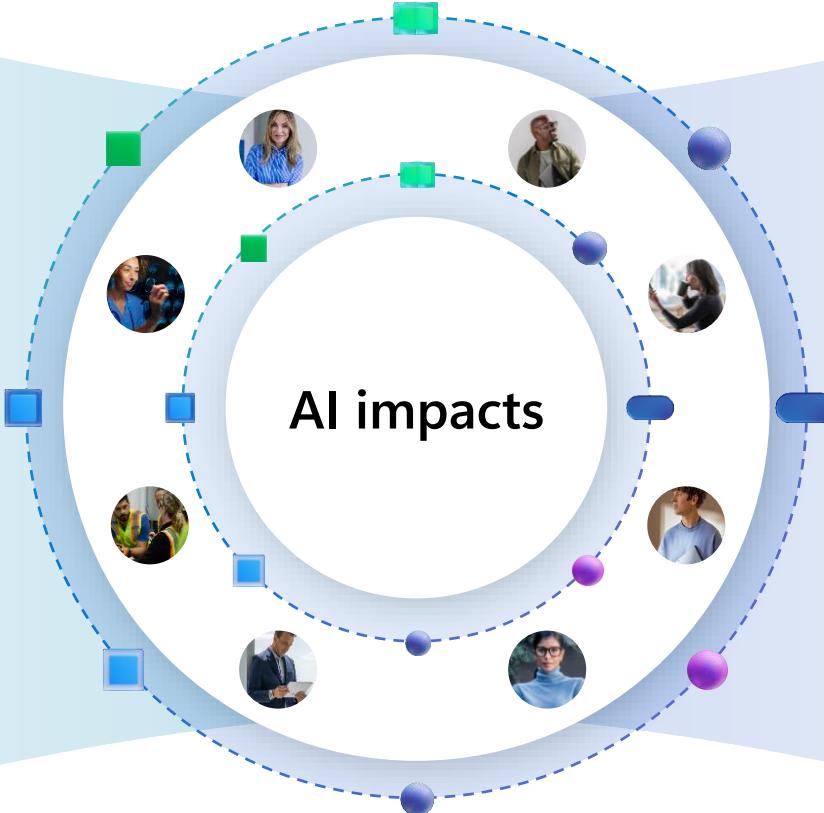
- AI tools for employee productivity
- Personalized consumer engagement
- Efficiency and profitability boosters

## Manufacturing

- AI copilots for product and maintenance support
- Innovation in R&D and design
- Engineering and back-office efficiency

## Telecom

- AI-enhanced contact centers
- Productivity tools for field and office
- Knowledge and process automation



<https://aiusecaseexplorer.microsoft.com/>

# Win Formula for Partners : Innovate with Azure AI Apps and Agents

## Customer Outcome

Build, run, and manage powerful AI apps and agents securely and cost-effectively at scale

## Partner Outcome

Help drive revenue growth, profitability, and margins by enriching your services with differentiated solutions—supported by Microsoft's full tech stack and Win Formula.

### Differentiated Partner Capabilities

Differentiate your value to customers by leveraging Azure's tech and GTM differentiation

Align customer goals to AI use cases and deploy repeatable solutions with security foundations

Drive Azure's full stack to engage with customers across their business and IT functions, also enabling AI CoEs

### Win Formula aligning to MCEM

#### 1 Listen & consult

##### Use-case identification and AI on Azure conversation

Know your customer: Target customers from [SPARK Propensity](#)

Understand priorities and establish base value:

- [AI Executive Conversation](#)
- Azure Conversation Guide (coming soon)
- Discovery questions (in this playbook)
- Use-cases (in this playbook)

#### 2 Inspire & design

##### Envisioning and solution design

Identify detailed requirements and pitch the play:

- [Biz & Capability Envisioning Workshops](#)
- [Solution Play L100 and L200 pitch decks](#)
- [Agentic DevOps L100 and L200 deck](#)
- [Product pitch decks](#)
- [Industry Sales Kit](#)

#### 3 Empower & achieve

##### Drive POC and customer skilling

- Accelerate proof-of-concept with [Gold Standard Accelerators](#) and [Azure AI Foundry Templates](#)

#### 4 Realize value

##### Deploy & Drive Adoption

Leverage [landing zone](#) and [deployment templates](#)

Validate solution against [Well Architected AI Assessment](#)

#### 5 Manage & optimize

##### Drive Optimization and Expansion

Leverage [Azure Essentials for resiliency](#)

Establish [AI Center of Excellence](#) at customer account

Engage BDMs for next AI use-case

# Key Actions – Complete ASAP!

*Use the Win Formula for sales success. But the below actions may increase your chances!*

Register for skilling

**aka.ms/partnerskilling**

Get certified (tech roles)

**aka.ms/certweek**

Get Specialized\*

**aka.ms/AIAppsSpec**

\*Azure Specializations are achieved by partner organizations (not individuals). Either of the following will enable Azure Accelerate benefits:

- [Build AI Apps on Azure](#) (Preferred, since it establishes credibility to deliver full stack implementations)
- [AI Platform on Azure](#)
- [Developer Productivity on Azure](#)



Drive Apps that  
Matter and establish  
AI Foundry as the  
platform of choice





# Apps that Matter

Build AI Apps that drive Frontier Transformation



## Business Impact

Delivers clear, measurable value by solving a real problem aligned with user needs and organizational goals:  
high-value use cases

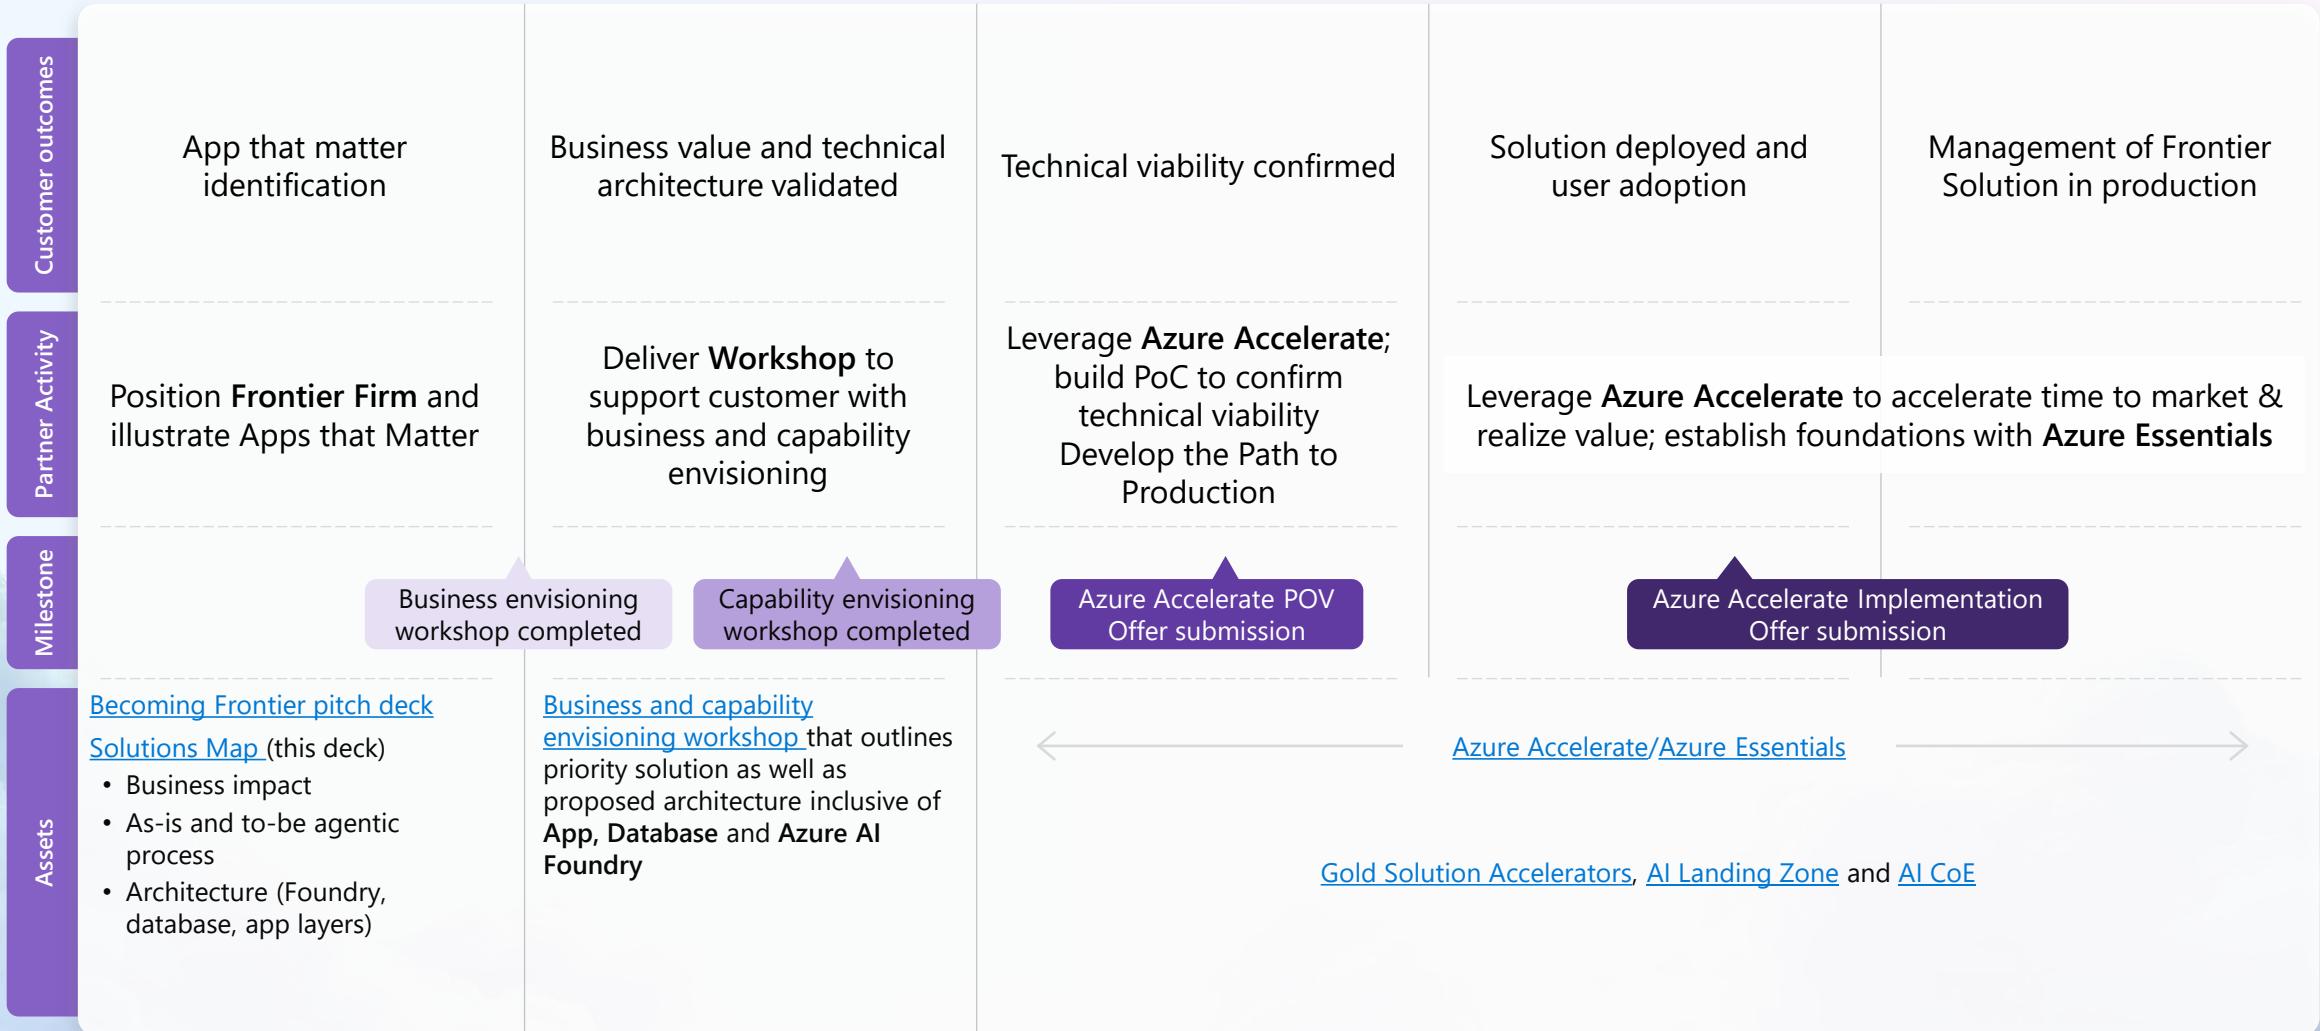
## Innovative

Leverages cutting-edge agentic technologies and real time data to solve complex problems in ways that are innovative, adaptive, and often predictive or autonomous

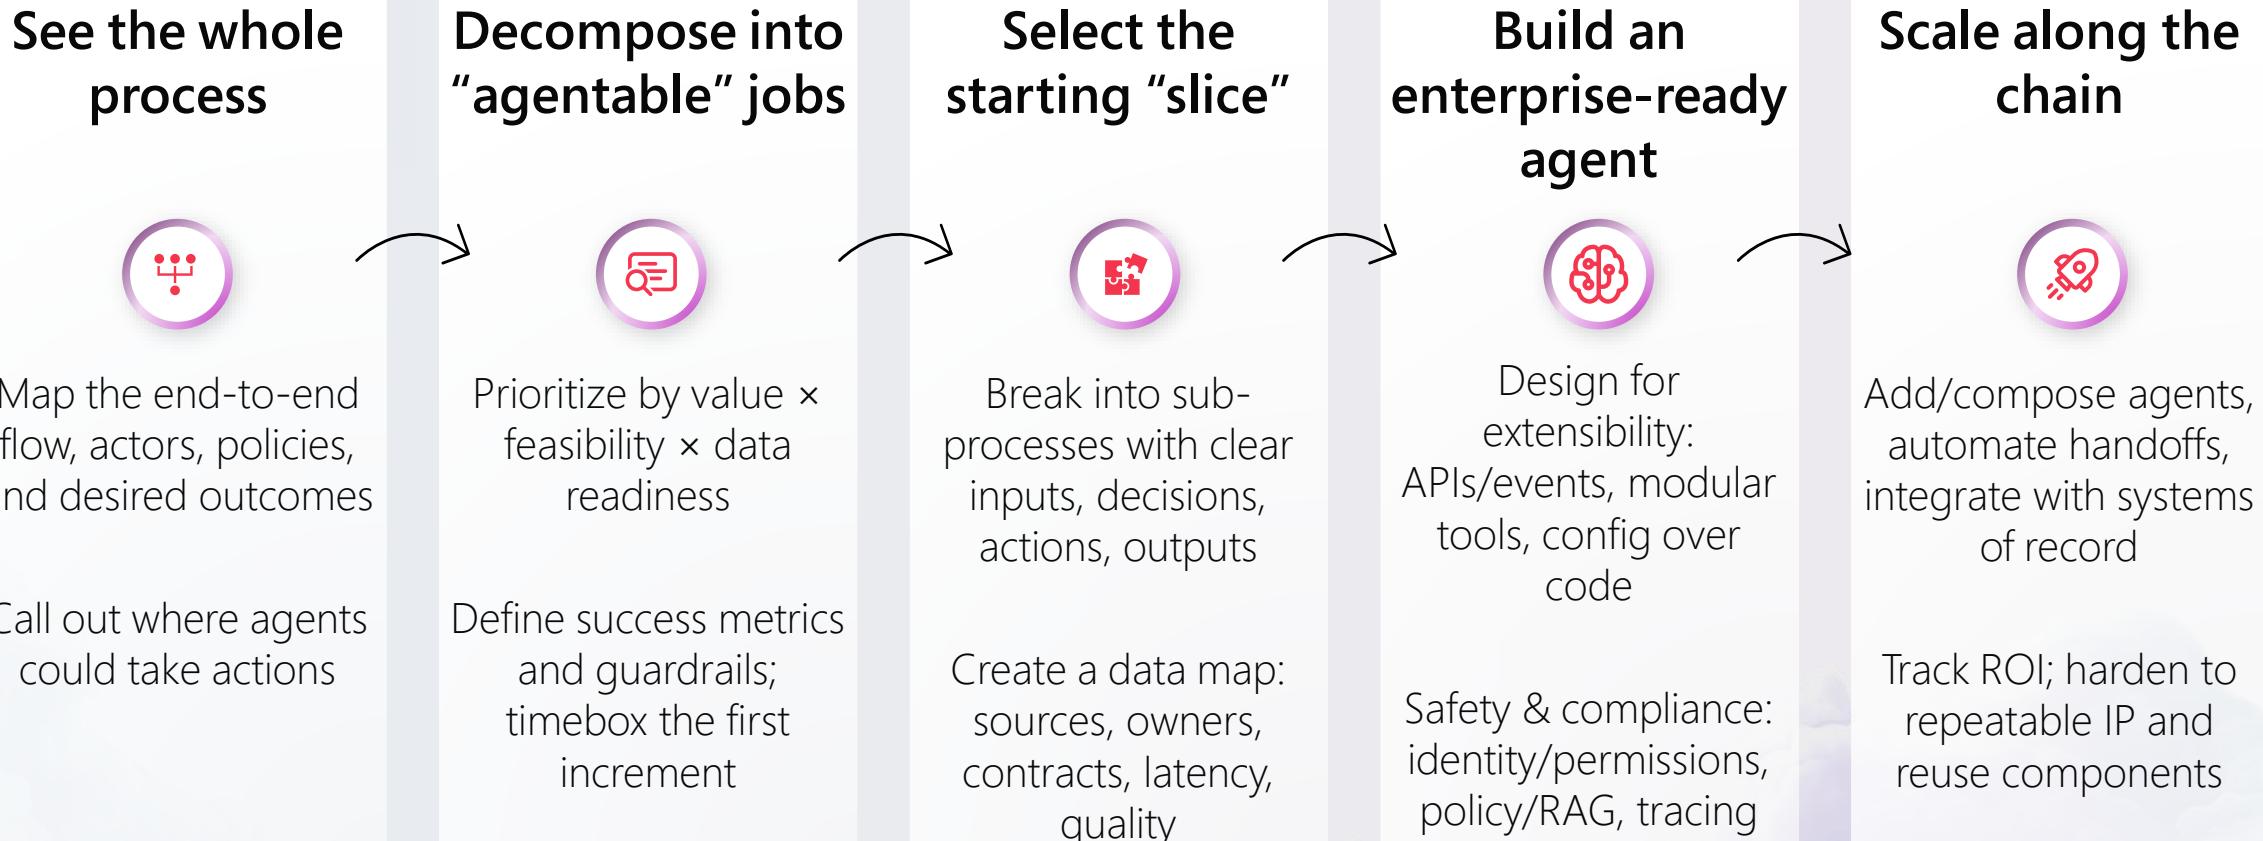
## Enterprise-grade

Production-grade systems that are governed, secure, trusted, operationally stable, business-aligned, continuously optimized

# Ideation to delivery for Apps that Matter



# Guardrails on how to work with the suggested processes



**Ground rules:** Agentic processes are a combination of multiple software solutions



# Using Envisioning Workshops to Drive Apps that Matter

Workshops to support scaling agentic automation leading to customer upskilling and a fully tailored agentic solution ready to scale



## Workshop 1: Agentic business envisioning

- Introduction to (*multi-*) *agent solutions* and the Microsoft AI Agent Platform
- Present *Apps that Matters* and reference examples including impact
- Dive deep into customer specific key processes
- Identify *Apps that Matter* to the customer

2–4hr workshop

[Link](#)



Delivery

## Workshop 2: Agentic capabilities for scale

- Whiteboarding solutioning to translate solution outcomes into an agentic process
- Dive deep on core elements of enterprise readiness (platform, security, responsible AI,...)
- Rapid prototyping to validate feasibility of uncertain parts of the solution
- Alignment on scaling path (via Azure Accelerate)

4–8hr workshop

[Link](#)

Case study  
after commitment

# Apps that Matter – Solution maps

## Cross-industry solutions

- 01 Order-to-cash process
- 02 Procurement & sourcing
- 03 Contact center operations
- 04 Contract lifecycle management
- 05 E2E supply chain operations

## FSI

- 01 Loan processing
- 02 Mortgage processing
- 03 Fraud analysis & AML mgmt.
- 04 Claims processes
- 05 Market research and analytics

## RCG

- 01 Interactive brand and product agent
- 02 Content generation and operations
- 03 Agentic commerce (shopping through 3<sup>rd</sup>-party agent)
- 04 Inventory planning and trend forecasting
- 05 R&D and product design

## Telco & Media

- 01 Pred. assurance & outage recovery
- 02 Field operations
- 03 Customer-facing engagement
- 04 Agentic advertising
- 05 Creative & content co-creation

## Manufacturing & Mobility

- 01 Agentic co-innovation
- 02 Predictive maintenance & asset orchestration
- 03 Knowledge management and preservation
- 04 Factory ops with real-time data
- 05 Engineering-to-order/requirements

Further industries to be added



# FSI Apps that Matter

## Overall industry specifics:

- Augmentation-first due to risk/compliance: Human-in-the-loop by design; automate only low-risk, reversible steps; require explainability, audit trails, and approvals
- Ecosystem over custom build: Greater reliance on specialized third-party platforms (KYC/AML, fraud, core banking); agents compose and trigger via APIs rather than replace

Key metrics relevant to the industry

1	<b>Loan application processing</b>	Collecting applicant data and documents, running KYC/AML and credit scoring, and routing to underwriting for an approve/decline with terms	
2	<b>Mortgage processing</b>	Collecting borrower and property data, verifying identity/income, performing credit/affordability and collateral checks to underwriting for approve/conditions/decline with closing disclosures	
3	<b>Fraud &amp; AML management</b>	Monitoring transactions and accounts in real time with rules + ML, enriching with device/behavioural signals, and triaging alerts to cases for timely block/release decisions	
4	<b>Claims processing</b>	Taking FNOL to settlement by checking coverage, assessing liability, reviewing documentation, coordinating repair/replacement, and pursuing subrogation where applicable	
5	<b>Market research &amp; analytics</b>	Analyze market data and news, economic indicators, earnings calls extracting key insights, building financial models, identify investment opportunities, and creating market and investment reports including tailored insights for clients	

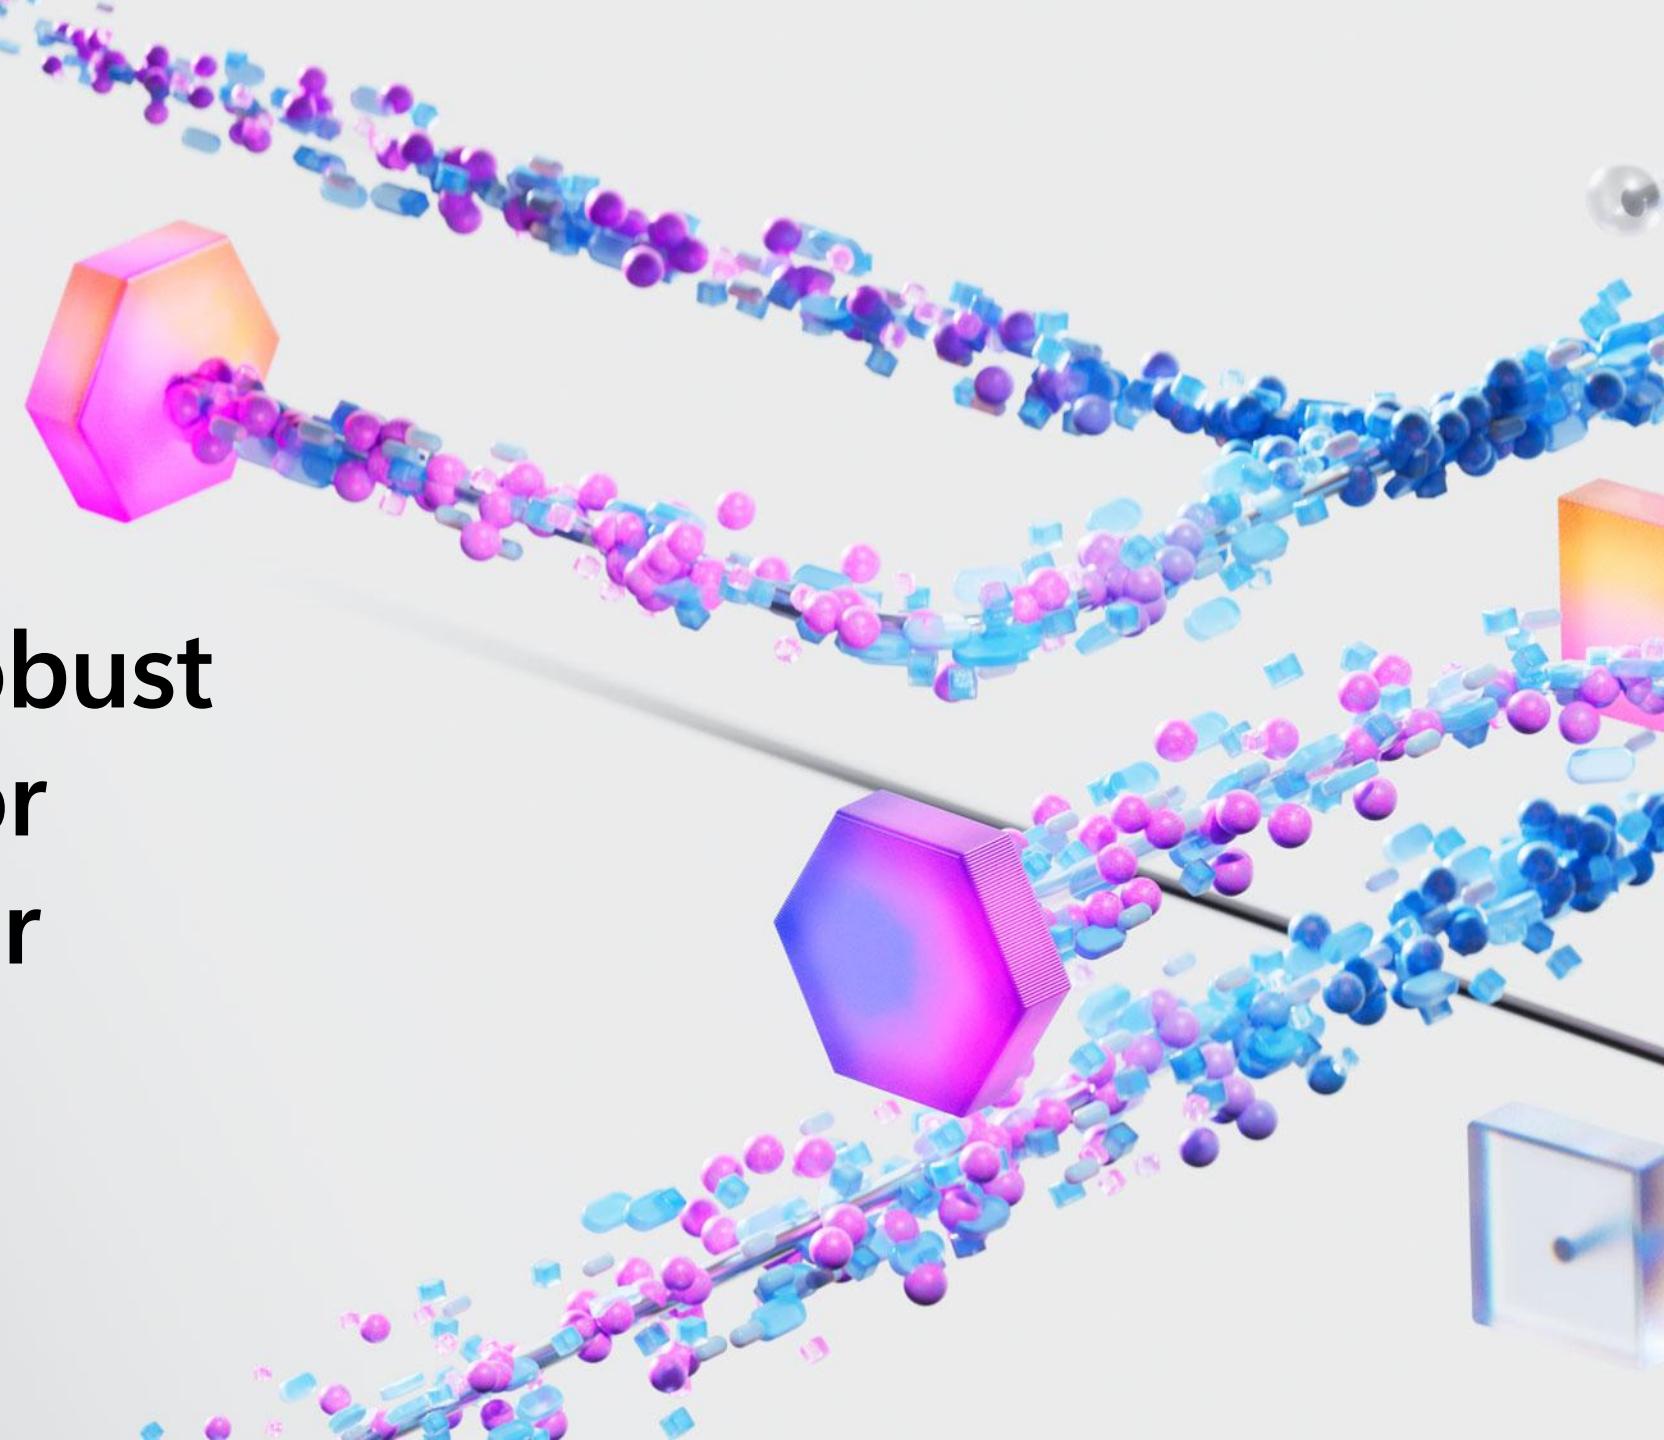
Banking focus

Insurance focus

Capital Markets

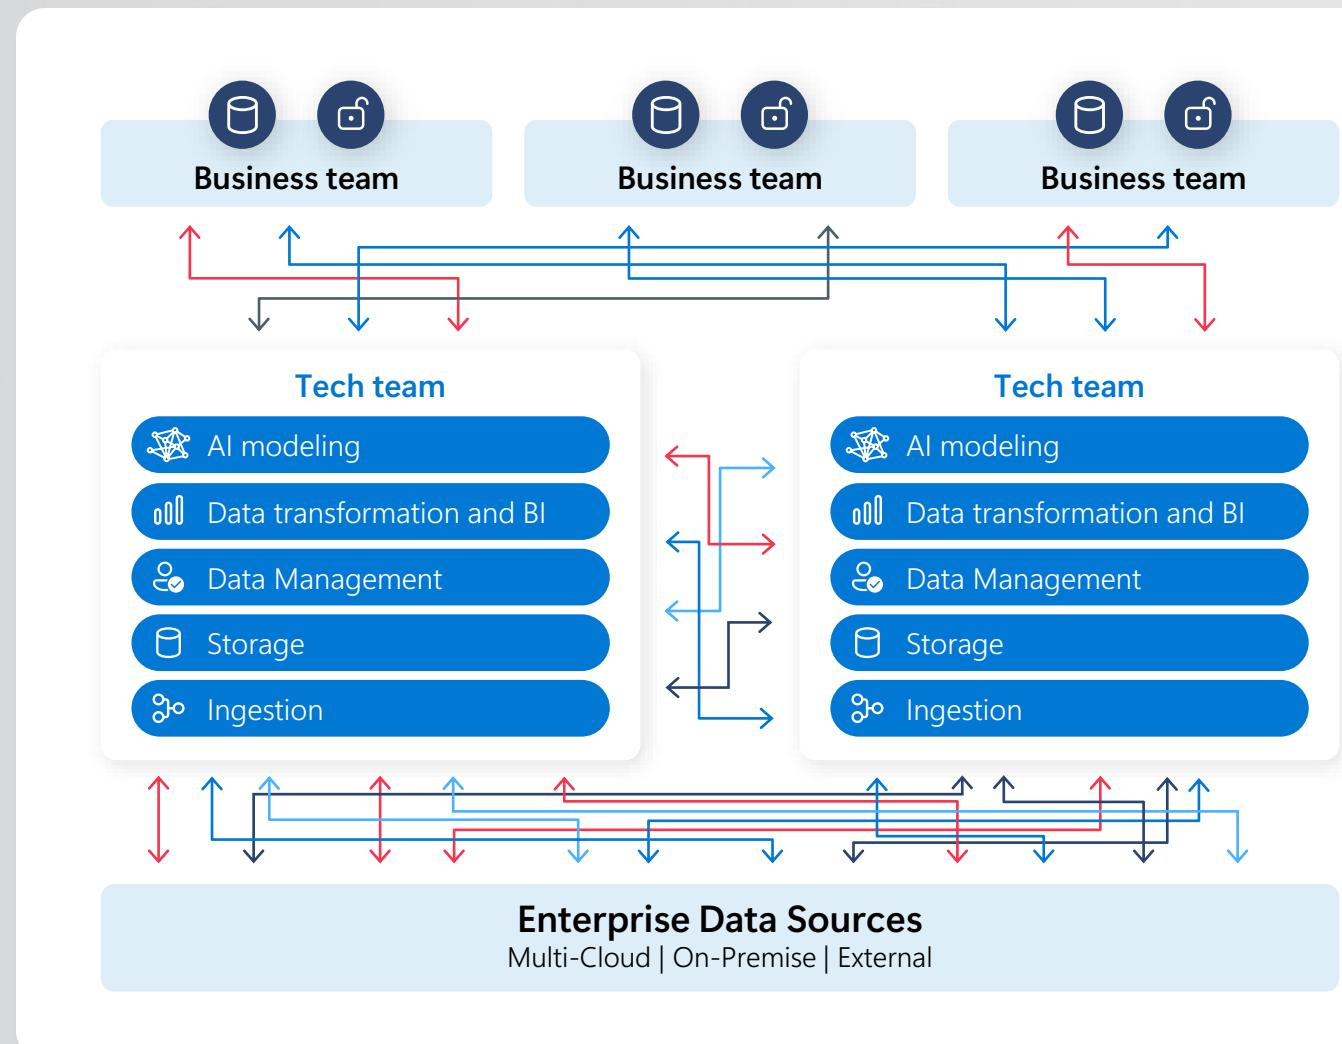
Sub-section

## **Establishing a robust data platform for Apps that Matter**



# The starting line

A complex, organically evolved data estate



- 1 Data copies and infrastructure inefficiencies
- 2 Limited interoperability
- 3 Data exposure risks
- 4 Cannot activate data for AI and innovation

# Accelerate data value creation with a modern data and analytics platform



**Data copies and infrastructure inefficiencies**



**Limited interoperability**



**Data exposure risks**



**Cannot activate data for AI and innovation**



Single, org-wide data lake that facilitates responsible data discovery and use



Unified source of data that integrates with databases for smarter decisions and scalability



End-to-end visibility and governance to reduce risk and maintain trust



Unified data platform optimized for AI scenarios

# Azure Databases solve real problems

Customers turn to Azure managed databases to gain scale, performance, agility, and simplicity that enables their teams to turn AI ambitions into business results.

Deliver real-time, always-on performance

Simplify data management and free up IT and dev teams

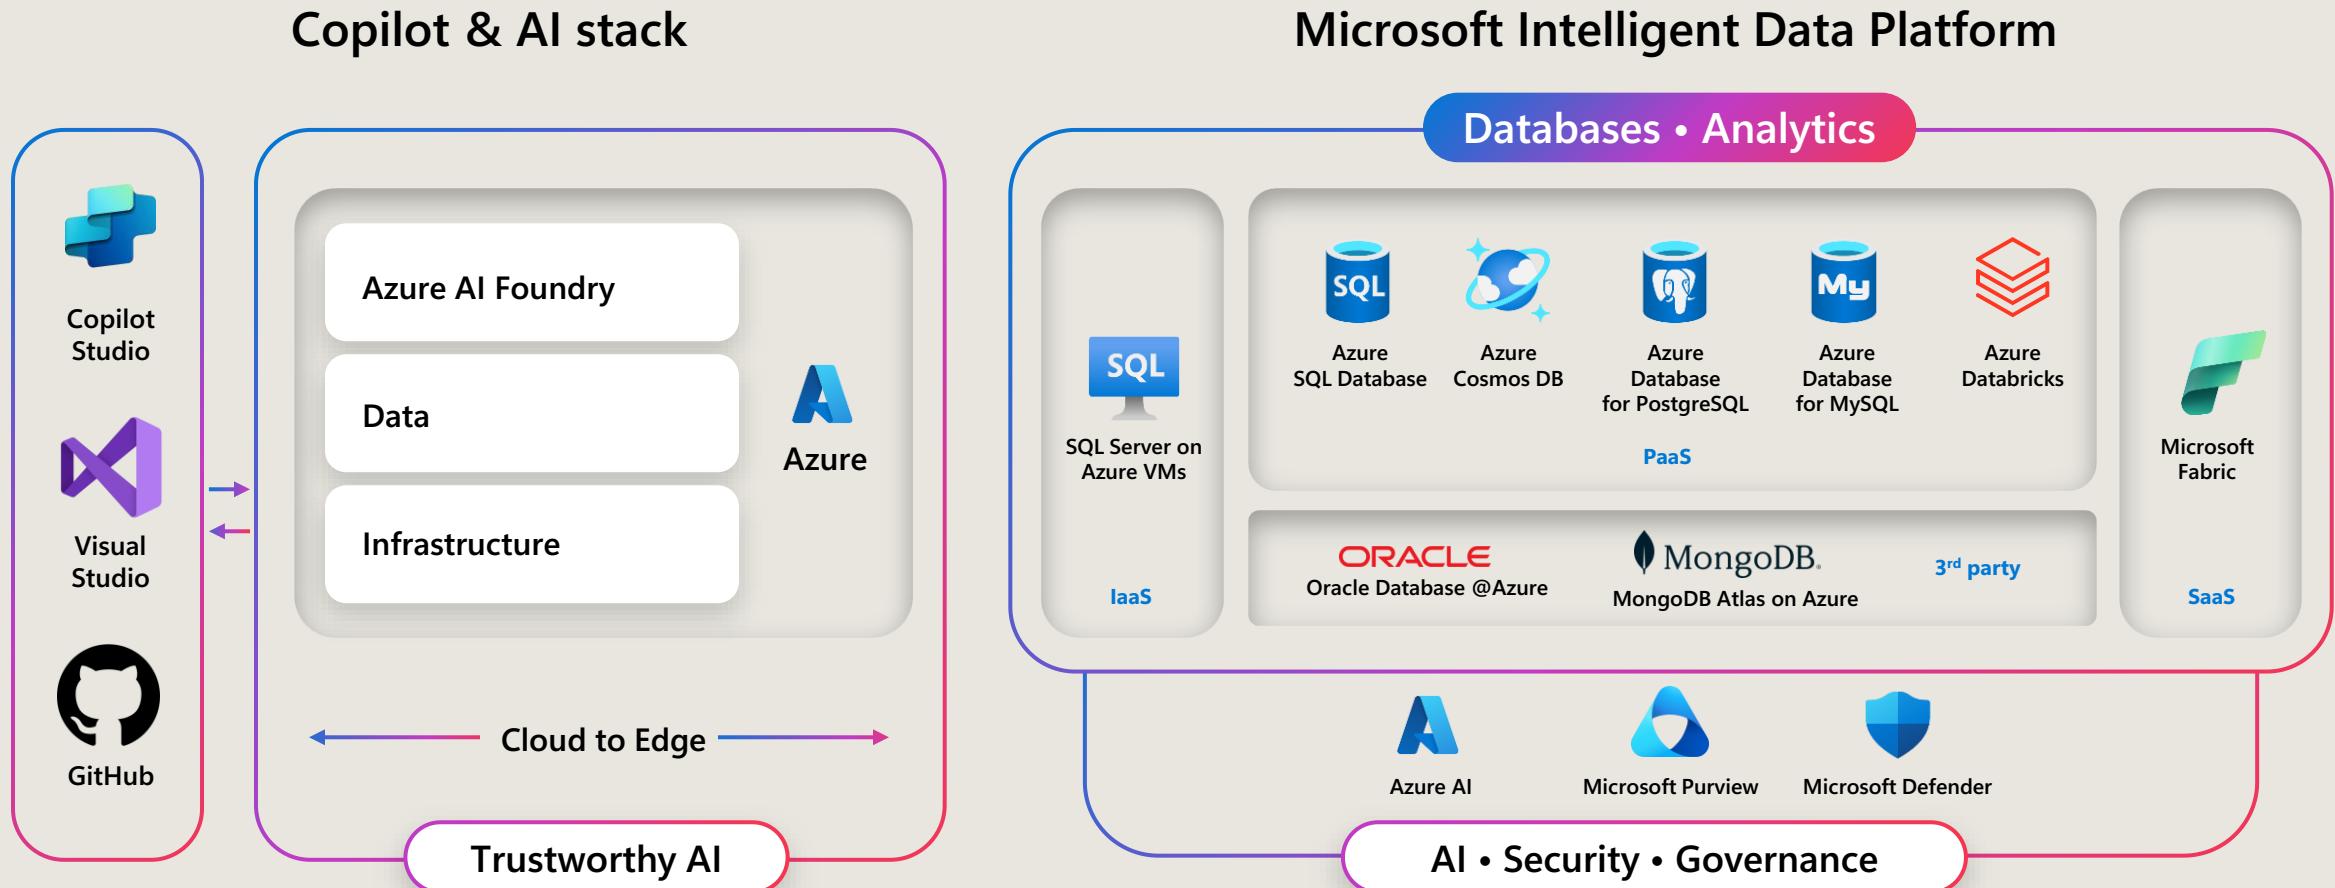
Respond with agility to changes, and adopt the newest tech

Improve security and compliance



# The finish line

A unified intelligent data platform built in the cloud and optimized for AI



# Microsoft Databases portfolio

Microsoft databases empower every developer and business to innovate without constraint and achieve the best time to value.

On-Premise



SQL  
Server

PaaS



Azure  
SQL Database



Azure  
Cosmos DB



Azure Database  
for PostgreSQL



Azure Database  
for MySQL

SaaS



SQL database in Fabric  
Cosmos DB in Fabric

3<sup>rd</sup> party

ORACLE

Oracle Database  
@Azure

MongoDB®

MongoDB Atlas  
on Azure

# Is your customer's database optimized for AI?



## Data at AI scale

Support workloads **of any size**

**Flexibly adapt** to changes and new requirements

Elasticity to **support dynamic workloads**



## Data in real-time

Provide AI models and analytics platforms with **immediate access to data**

**Low-latency reads and writes** to apps and agents

**Make data available** to different teams, apps, models, and processes



## Fast and efficient vector support

Ability to **natively store and query vectors**

Advanced information discovery with **vector search, reranking, and semantic search**

Highly-efficient and **cost-effective approximate nearest neighbor (ANN)** support



## Powerful integrations

Built-in **support for AI models and tools**

Enable advanced analytics through **real-time access for OLAP systems**

**Easily integrate** with Azure AI Foundry, Azure Databricks, and Microsoft Fabric

# Database Security

**End-to-End Encryption:** Secures data during storage and transfer, blocking unauthorized access.

**Role-Based Access Control (RBAC):** Limits permissions based on user roles to prevent overreach.

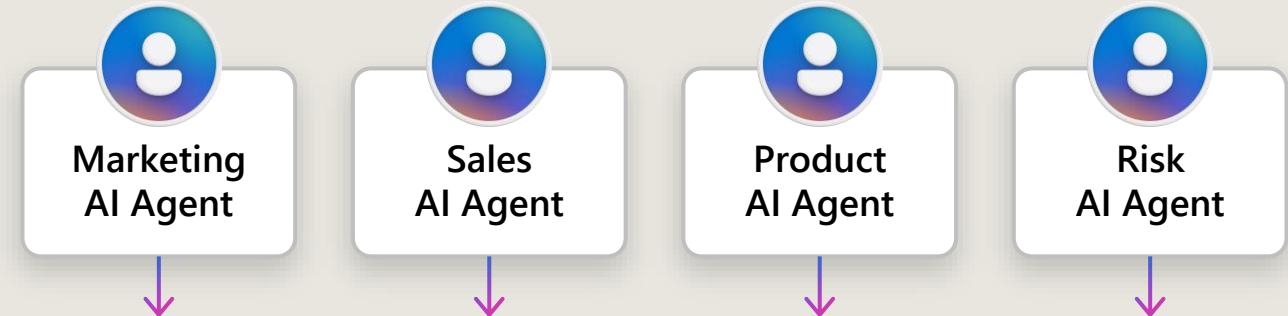
**Comprehensive Audit Trails:** Tracks all queries for complete data access visibility.

**AI Query Validation:** Restricts questionable queries before accessing sensitive data.

**Regulatory Compliance:** Adheres to GDPR, CCPA, and HIPAA standards for data protection.

**Key Message:** A database architecture is designed to secure and protect customer data -ensuring that AI-driven agent queries remain controlled and compliant.

AI Agent represent any systems or automated processes generating queries on customer data.



**Database**  
Database Security Layers  
(AI Query filters)

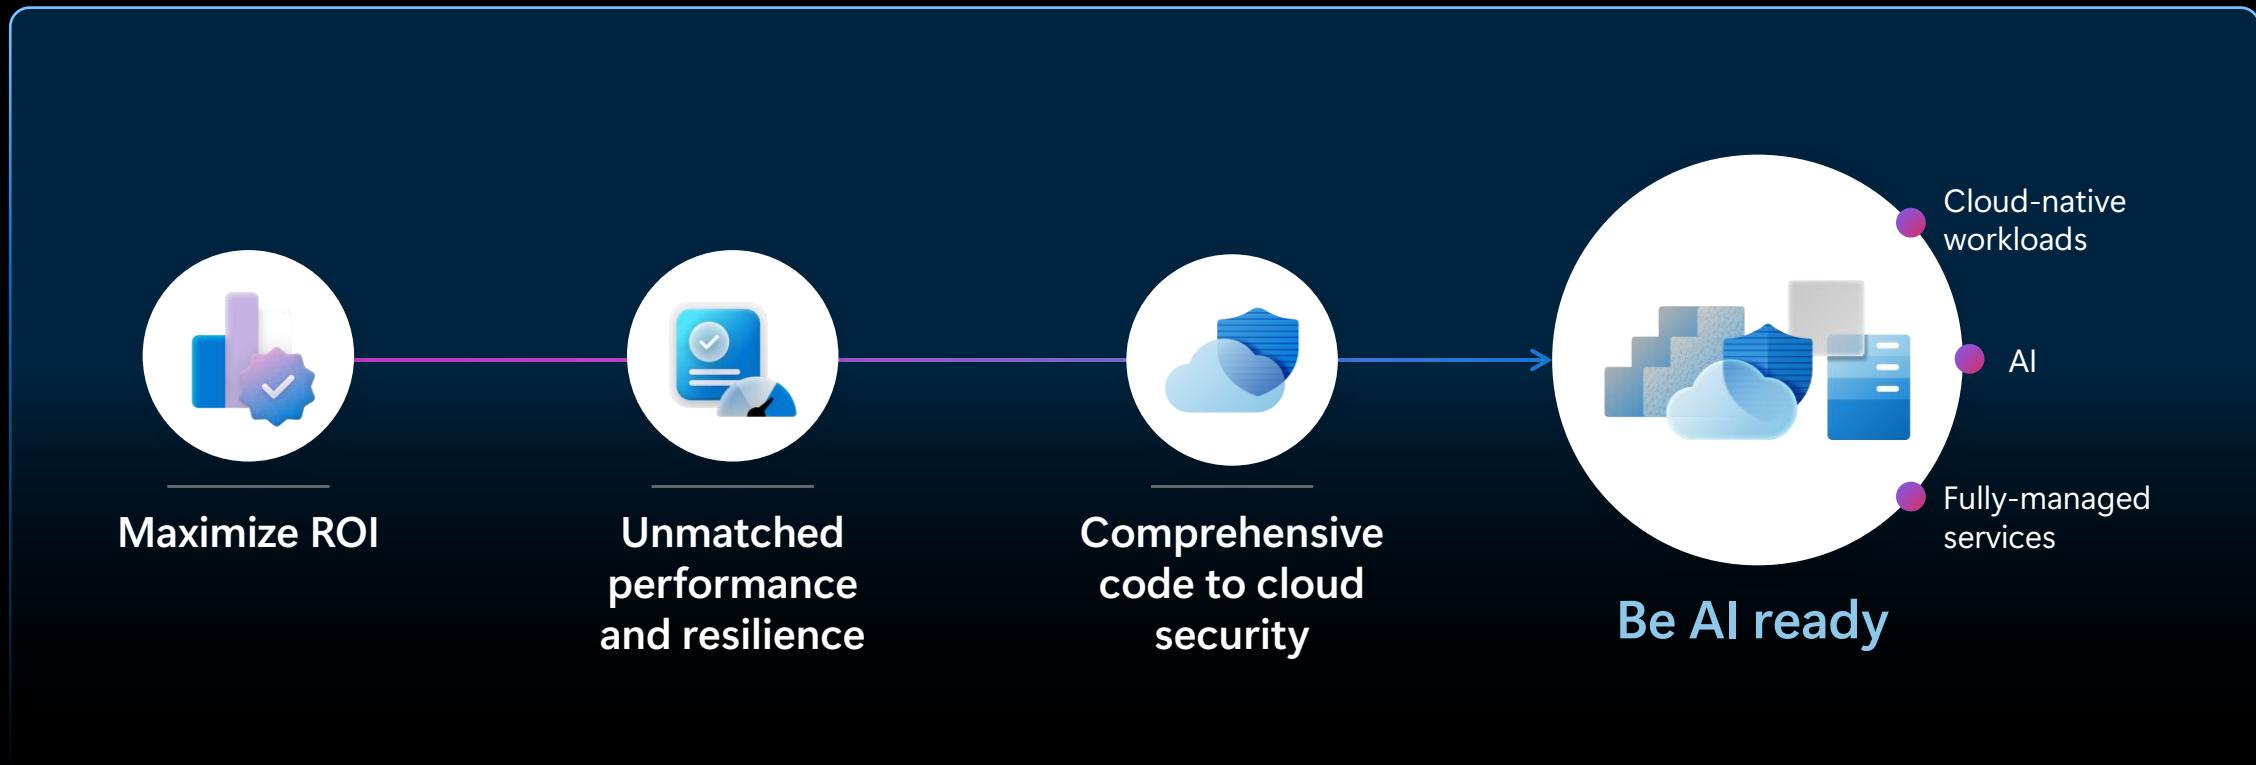
- Role-Based Access
- Data Encryption
- Logging & Auditing
- Row Level Security



- Customer Data
- Sensitive Information

- **Compliance (GDPR, HIPAA, Etc)**
  - Data Encrypted at rest & In Transit
  - Controlled Access
  - Data Masking
  - Secure Back-up & Monitoring

# The Microsoft cloud is the path to innovation



# Discovery Questions

- ✓ Identify strategic use cases and technical patterns for transforming business processes with AI
- ✓ Leverage these discoveries to implement AI solutions tailored to specific customer needs and objectives

## Business Considerations & Use Case

- In simple language, what problem are you trying to solve?
- How are you doing this business process today? What improvement would you like to see (acceleration, customer experience, cost reduction, etc.)?
- What would this solution enable your business to do? How will you measure success?
- Who are the users of the system? What is your desired user experience/How will the users interact with this solution?

## App Architecture

- How will AI address the problem? How do these rank in importance of achieving your business outcomes: **Availability, latency, security, geo distribution, cost, accuracy, speed to market?** *These have implications on the solution design.*
- **Are you planning to build a new AI app, or extend an existing app to add AI capability?**
  - If the customer is extending an existing app, where is the app currently hosted (Microsoft/Azure or outside)? How much customization will be required for the solution?
- What other applications would connect to this system? 3<sup>rd</sup> party or SaaS?
- Are you planning to co-locate your data with your AI solution at scale?
- How will the business application integrate with existing systems and ensure scalability for future needs?

## AI Model

- Have you experimented with model APIs? Are you planning to use a specific AI model? Why choose that model? (response time, accuracy, scale)
  - Are you considering building your own AI model? Are you training a new model? How do you plan to train?
  - **What are you using to keep up with changes in data (Data Drift), model**
- (LLMOPs), apps (DevOps) or outcomes (Model evaluation & monitoring)?**
- How will you ensure the AI system is fair and unbiased?
  - What measures will you take to ensure the AI system is transparent and explainable?

## Data Considerations

- What are the sources & formats of the data to be utilized for this project? Is there a governance framework in place for the datasets, managing both data quality and access?
  - How large is data and how fast does it change?
- **Will the solution involve any form of enterprise or non-public data that is stored in Enterprise Data Lake?**
  - How will you address data privacy and security concerns with the AI system?

## Infra and Security

- What is your strategy to scale your AI solution, and what outcomes are you looking to achieve?
  - How are you planning your foundational infrastructure to support your AI solution at scale? Do you have an AI landing zone?
- Have you reviewed your network architecture for your AI solution at scale?
  - What is your security model for your AI solution?
  - How are you thinking about governance and controls for your AI solution?
  - How are you thinking about cost management and ROI measurement?

# How do I navigate the DB conversation with the critical stakeholders?

## 1 CTO VP of Engineering

Set the stage for the customer to understand that to achieve the full potential of AI, it needs to build using modern Azure Databases and Modernize its legacy Databases.

Identify leadership priorities and key stakeholders in charge of landing them.

## 2 Engineering Director Enterprise Architect

Build trust, confidence on Azure DBs' features, TCO, use cases, security, tooling, accelerators, offers, programs, activities, customer stories.

Identify teams that need to be further engaged to land a detailed proposal, PoC

## 3 Dev. Team Lead Dev. Team

Build trust and enable Technical decision Makers, Devs. Finalize architecture design, fine tune TCO, land the adequate programs and investments. Execute PoC

Turn teams into advocates.  
Land a successful PoC. Win technical decision.

# Strategic Drivers Behind Database Transformations

Barriers to  
Innovation

From:

Complex Data integration layers impose latency and complexity, add cost.

Different DBs technologies slow down App development, demand for more skills.

Traditional DBs lack critical AI capabilities for AI Apps. Slow Dev Cycles. Complexity.

File System based Docs Data lack security and performance objectives for AI Apps.

External AI, Search can increase latency and raise performance, security challenges.

Your business process performance is as good as your worst DB performance.

Security, compliance harder to implement across a too diverse data stack.

Insights dependent on ETL, batch processes. Real-time intelligence harder.

To:

Simple, performant, standard and scalable Databases Speed layer to fuel your Apps.

Accelerate Dev through standardized Data platform. Less skills needs.

Modern, Cloud based, AI ready, Future Proof Databases.

App relevant Doc. data standardized in NoSQL real time, scalable DBs.

Run AI next to Data, reducing latency, risk, increasing simplicity, performance.

Consistent performance guarantees across all Apps and core systems.

Standard and automated Security, Compliance and audit mechanisms.

Leverage Fabric Databases. Merge Operational and Analytical.

# Six Key Customer Priorities Driving Database Transformation

1

## Barriers to Innovation

With AI becoming mainstream, CxOs are eager to harness these technologies, which raises concerns about data readiness and integration. AI's potential can offer game-changing insights and automation, but only if the underlying data architecture supports it. A common CxO question is: "Do we have an AI-ready data foundation? CxOs are also cautious about dependence on any single database vendor or technology. They worry that once a system is in place, it could be very costly or difficult to replace.

3

## Longer Time to Market

CxOs commonly cite data silos and fragmentation as obstacles. Data scattered across disparate systems makes it hard to get a unified view of the business and requires many different technical capabilities in the workforce.

5

## Inefficient Performance and Scalability

Executives worry whether their data infrastructure can handle ever-growing data volumes and user demand without performance degradation. With an explosion in data, CxOs prioritize scalable database solutions that maintain fast response times as the business and its datasets expand.

2

## High Operating Costs

The cost of data platforms and projects is a key concern at the CxO level (especially for CFOs and CIOs). Leaders must ensure a strong return on investment for expensive database technologies.

4

## Inconsistent Security Posture

CxOs must be confident that if a major incident occurs: server failures, cyberattack, natural disaster, or even human error, the company's data is safe and can be quickly restored. They plan for the worst-case scenarios: robust backup, restore, and continuity plans. Managing cybersecurity risk is a top focus for 69% of CIOs, as well as ensuring compliance with data privacy laws (GDPR, CCPA, HIPAA, etc.).

6

## Real-Time Data Processing

CxOs increasingly demand real-time data capabilities. They are concerned that their data architecture only supports slow/batch processing. In the age of AI, modern use cases, from live customer personalization to fraud detection require streaming and real-time data processing.

1

# Barriers to Innovation

**Data silos and fragmentation are obstacles. Data scattered across disparate systems makes it hard to get a unified view of the business and requires many different technical capabilities in the workforce.**

## Questions to Ask:

- Do you have data in different systems that don't talk to each other, and what challenges does that create for your team?
- How do you currently combine data from various sources, and are there pain points in that integration process?
- How can modernizing our databases to PaaS accelerate AI initiatives and innovation in our business?

## Azure Databases Talking points

- **AI-Ready Data Platform**: Adopting a PaaS database creates a unified, cloud-based data foundation that is built to infuse AI and analytics through "intelligent cloud databases".
- **Unified Data Access**: Cloud databases consolidate structured and unstructured data into a single platform, giving leaders a complete, real-time view of the business and enabling seamless cross-functional collaboration.
- **Multi-Model Flexibility**: A unified architecture supports diverse data types - SQL, NoSQL, analytics, reducing the need for specialized platforms and skills, fostering broader team collaboration.
- **Operational Simplicity**: Replacing fragmented on-prem systems with cloud platforms reduces complexity and maintenance overhead. IT teams can focus on delivering insights, while employees access data through familiar tools, boosting productivity and adoption.

2

## High Operating Costs

### Look out for:

- Legacy Databases and End of Support events
- Large Database management teams, stretched with “upkeep” tasks
- Frequent fire fighting to overcome availability and capacity issues
- On-prem workloads

Migrating to Azure’s managed database services can significantly lower costs and improve operational efficiency. Leaders must ensure a strong return on investment for expensive database technologies.

### Questions to Ask:

- What percentage of your IT budget is consumed by maintaining DBs (hardware refreshes, licensing, support) versus investing in new initiatives?
- Do you often over-provision DB capacity to handle peak loads, and therefore pay for a lot of idle resources during normal times?
- How much manual effort does your team spend on routine DB tasks like patching, backups, and scaling? If those were automated, could your staff focus on higher-value projects?

### Azure Databases Talking points

- **Lower Total Cost of Ownership (TCO):** Eliminate large capital expenditures on servers and storage and shift to a pay-as-you-go model. Stop paying for unused capacity. More predictable costs, and easier to align with business priorities.
- **Increased IT and Staff Efficiency:** Less time on upkeep (patching, upgrades, backups, and monitoring) that can be reallocated to strategic work. A smaller team can effectively oversee a larger environment: doing more with less. Deliver new projects faster and better support the business.
- **Optimized Pricing and Discounts:** Azure Hybrid Benefit allows you to reuse existing on-premises SQL Server licenses in the cloud, cutting license costs by up to 30-50%. For predictable workloads, you can purchase reserved capacity. The net effect is substantially improved cost-efficiency versus traditional models. Features like auto-pausing DBs when idle (for certain services) and serverless compute tiers ensure you aren’t paying a penny when your DBs aren’t in use.

### 3

## Longer Time to Market

#### Look out for:

- Long Development cycles
- Difficulties in going from Idea, PoC to production applications
- Lack of Dev team autonomy to rapidly have “ready to use” Database infrastructure

**Migrating to Azure’s managed database services can significantly accelerate Development Cycles. Leaders must ensure a strong return not only on investment for expensive database technologies but also in Productivity.**

#### Questions to Ask:

- Have you had initiatives to break down data silos and consolidate your Databases, and what obstacles have you encountered in those efforts?
- Did you know there are significant performance and scalability benefits of using PaaS databases for your AI and analytics workloads?
- What is the impact on cost and ROI if we move to PaaS databases for AI acceleration, and how can it improve our overall cost, management (TCO)?

#### Azure Databases Talking points

- **Real-Time Consistency:** Synchronized data across systems ensures everyone works from the same source of truth, eliminating delays and enabling faster, data-driven decisions at all levels.
- **Faster AI Development Cycles :** Developers and data scientists are freed to focus on AI innovation rather than provisioning and tuning databases. New environments can be spun up in minutes, enabling rapid prototyping of AI solutions. AI-enhanced projects come to market faster, which is crucial for staying ahead of competitors.
- **Seamless Integration :** Your data can flow into AI/ML pipelines with minimal friction. Such tight coupling accelerates the AI development lifecycle, data moves more quickly from collection to model training to insight. The result is a shorter path from raw data to intelligent action. Your team can iterate rapidly. The plumbing between operational data and AI tools is largely solved by the cloud platform.
- **Scalable Infrastructure :** As your AI goals expand (more data, more models, more users), the database infrastructure won’t be a bottleneck. Cloud databases scale on-demand – no long procurement cycles or capacity ceilings. Flexibility encourages innovation.

## 4

# Inconsistent Security Posture

### Look out for:

- Recent security breaches or issues
- More than 10 different Database technologies powering key business applications
- Legacy, EoL Database technologies

CxOs concerned about protecting sensitive data from cyberattacks and breaches, which can cause devastating financial and reputational damage. Managing cybersecurity risk is a top focus for 69% of CIOs.

### Questions to Ask:

- What data security or compliance challenges are you currently facing with your on-prem databases (e.g., meeting GDPR, HIPAA, or other regulations)?
- How do you control and audit access to sensitive data across your databases today? Are you confident only the right people have access, and can you easily track activities for compliance?
- How does your team handle security updates and threat monitoring for your databases? Do you worry about unpatched vulnerabilities or undetected suspicious activities in your data environment?

### Azure Databases Talking points

- **Built-in Data Encryption:** Azure PaaS databases encrypt data at rest and in transit by default (Transparent Data Encryption for storage, TLS for connections), so sensitive information is protected without any setup.
- **Unified Access Control:** Integration with Azure AD enables single sign-on and centralized role-based access control (RBAC) for database users. Permissions can be tightly managed in one place, ensuring only authorized personnel can access or modify data.
- **Advanced Threat Protection & Updates:** Continuous security monitoring and threat detection. OS and DB engine automatically regularly patched and updated engine.
- **Compliance & Governance Ready:** Moving to Azure PaaS gives you a governed data environment where security best practices and regulatory requirements are bolted in by design.

5

# Inefficient Performance and Scalability

## Look out for:

- Low resource usage in key Databases
- Capacity related outages or, pre-provisioning of capacity to respond to peaks
- Legacy, EoL Database technologies

CxOs depend on having Innovation, not only delivered faster and with a better TCO but also to guarantee it remains available, regardless of Scale and Performance. It's any CxO nightmare to invest on a new platform that cannot cope with its success.

### Questions to Ask:

- How long does it typically take to provision a new DB instance for a project or testing in your current environment? - This probes if slow setup times are hindering rapid development cycles.
- If your application usage doubled overnight, can your current DBs scale quickly enough (without significant rework or downtime) to handle the load?"
- Are your developers spending significant time on DB upkeep (patching, tuning, backups), and could that be slowing down delivery of new features?

### Azure Databases Talking points

- **On-Demand Scalability :** Azure PaaS DBs scale up or out on-the-fly to meet changing demands without complex re-architecture. Workloads can be easily scaled up or back as needed, so performance stays consistent at any load. This elasticity means your innovation isn't constrained by physical capacity.
- **Flexible, Cost-Efficient Environments :** Azure's pay-as-you-go model, special Dev/Test pricing and serverless options make it easy and affordable to create multiple DB environments on demand. Teams can spin up temporary DBs for QA, experimentation, or load testing without lengthy approval or big expenditures. This fosters an agile, experiment-friendly culture.
- **Dynamic Scalability :** Operational resiliency isn't just about surviving faults; it's also about smoothly handling unexpected spikes in demand. Azure PaaS DBs offer on-demand scalability that can mitigate outage risks from overload. If your application usage suddenly doubles, the DB can scale up without disrupting connections. Your DBs can be configured to autoscale within limits, and then resources dial back down to control cost. Resiliency against failures, and against success.

## 6

# Real-Time Data Processing

### Look out for:

- Complex Data pipelines
- Long, Time consuming and processing heavy batch ETL processes
- Heavy integration platforms

CxOs increasingly demand real-time data capabilities. They are concerned that their data architecture only supports slow/batch processing. In the age of AI, modern use cases, from live customer personalization to fraud detection require streaming and real-time data processing.

### Questions to Ask:

- How quickly can your organization react to new data as it's generated? For example, do you rely on nightly batch updates to see the latest metrics, or do you have a way to get instant visibility into ongoing transactions and events?
- Can your current databases handle high-velocity data streams (from customer interactions, IoT sensors, etc.) without lag or manual intervention to scale up? What happens to your system if data spikes unexpectedly – does performance suffer or require firefighting?

### Azure Databases Talking points

- **Ultra-Low Latency Operations :** Extremely fast read and write response times, which is critical for real-time processing. Applications can ingest data and retrieve results instantly, enabling use cases like live dashboards and rapid customer interactions. Compared to legacy on-prem systems (where queries might take seconds or more under load), a PaaS approach provides snappy, predictable performance even as data volumes grow.
- **Elastic Throughput Scaling :** On-demand, automatic scaling that can handle bursty data streams without manual intervention. If your company suddenly sees a surge in events like a traffic spike or a flood of IoT sensor readings. Capacity scales up during peak ingestion and scales back down during lulls, so real-time data keeps flowing smoothly without over-provisioning or hitting performance ceilings. PaaS ensures that as data streams grow, the platform seamlessly expands to keep up, which is vital for real-time insights or event processing pipelines that can't afford bottlenecks.
- **Fabric Integration :** Direct OLTP Database availability on Fabric, guaranteeing NoETL scenarios and Real time.



# Drive the full Azure tech stack – Product Highlights



# Build agents, your way

Powered by  
Azure AI Foundry



## Copilot Studio

Low-code managed environment for  
AI business process automation



## Visual Studio/GitHub

Pro-code environment  
for commercial AI solutions



## Azure AI Foundry

Comprehensive platform of capabilities, models, and tools to design,  
customize, and manage AI applications and agents



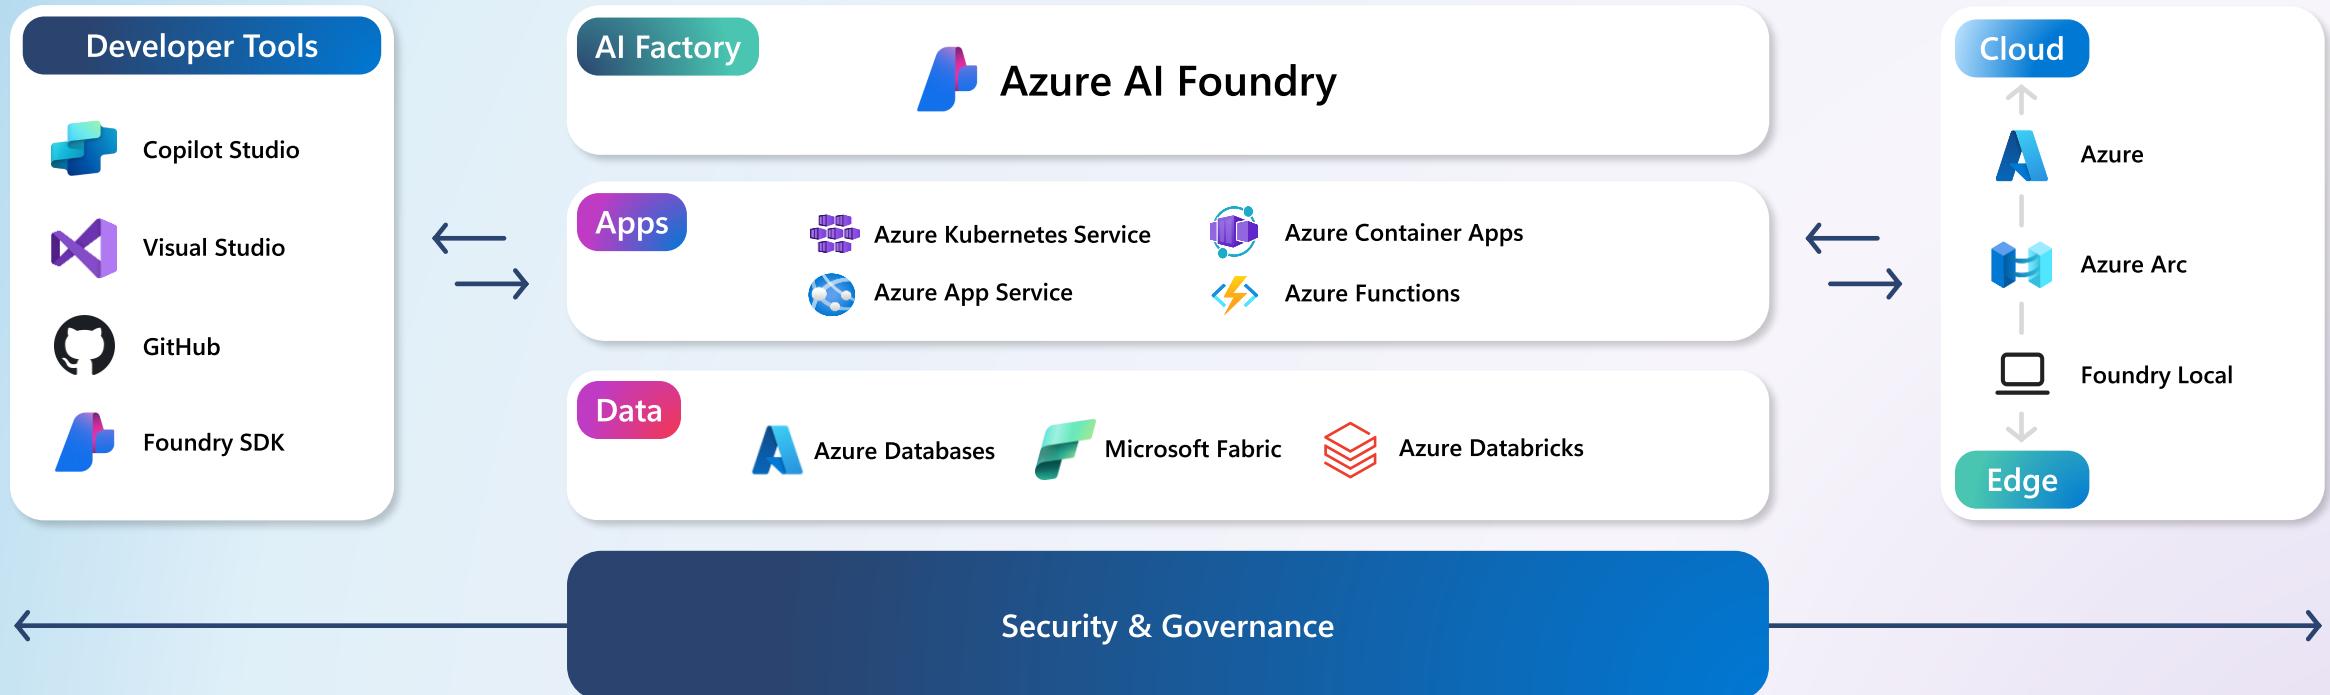
## Azure Databases

## Microsoft Fabric

## Azure Databricks

Complete data platform to accelerate AI innovation

# Azure AI apps and agents platform



Generally available



# Azure AI Foundry Agent Service

Securely design, deploy, and scale AI agents with ease

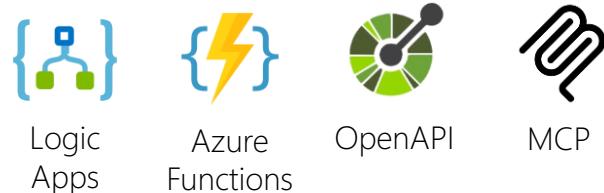
## Tools

Richest set of enterprise connectivity

## Knowledge



## Actions



## Choice

Model choice and flexibility with Foundry Models



### Azure OpenAI

GPT-4.1 series, o-series, GPT-4o



Meta models



Mistral models



Cohere models



DeepSeek models



Grok3 models from xAI



Hugging Face models

## Trust

Customer control over data, networking, and security

BYO-file storage

BYO-search index

BYO-virtual network

BYO-thread storage

Tracing/monitoring

Evaluation

# Azure AI Foundry Models

Offering 11,000+ frontier and open models

## Built-in safety

**Deployment Types:**  
Standard  
Provisioned Throughput  
Batch (50% off)

**Deployment Locations:**  
Global  
Data Zone (US and EU)  
28 Regions

**Sold by:**  
Microsoft

## Hosted and sold by Microsoft



OpenAI  
Model Family  
(available day 1)



DeepSeek  
latest



Mistral AI  
latest  
(coming soon)



Meta Llama  
latest  
(coming soon)



xAI  
Grok  
(coming soon)



Black Forest  
Labs  
(coming soon)



Microsoft Phi  
Family

## Other Foundry Models



Hugging Face



NVIDIA  
NIMS



Cohere

## Specialized & Industry models



**Deployment Type:**  
Standard  
Managed Compute

**Deployment Locations:**  
Global  
7 Regions

**Sold by:**  
Partner

# Guiding questions to choose the right model

What tasks are you trying to accomplish?

**NLP:** consider GPT-4, GPT-3.5, Copilot, Llama3, Mistral as your starting point

**Audio:** For Text to speech consider Open AI Whisper

**Computer Vision:** OpenAI DALL-E, Phi-3 Vision, Stability AI Stable Diffusion

**Multimodal:** OpenAI GPT-4o

Do you need high precision in a specialized domain?

- Your use case requires fine-tuning to match the precision and domain knowledge required: we recommend you leverage models like Llama, Phi

- If high performance is required on general purpose knowledge across diverse fields: most likely big proprietary LLMs will be your best bet.

- Otherwise, use either proprietary or open-source. This includes cases when extreme specialization of the model isn't required, since grounding knowledge may be achieved via RAGs.

What are your inference speed requirements?

- Fast response - Try smaller models (7b, or quantized) that fit into consumer hardware. Examples are Microsoft Phi3 and LLaMa3 8b. ChatGPT3.5 Turbo is quite fast as well.

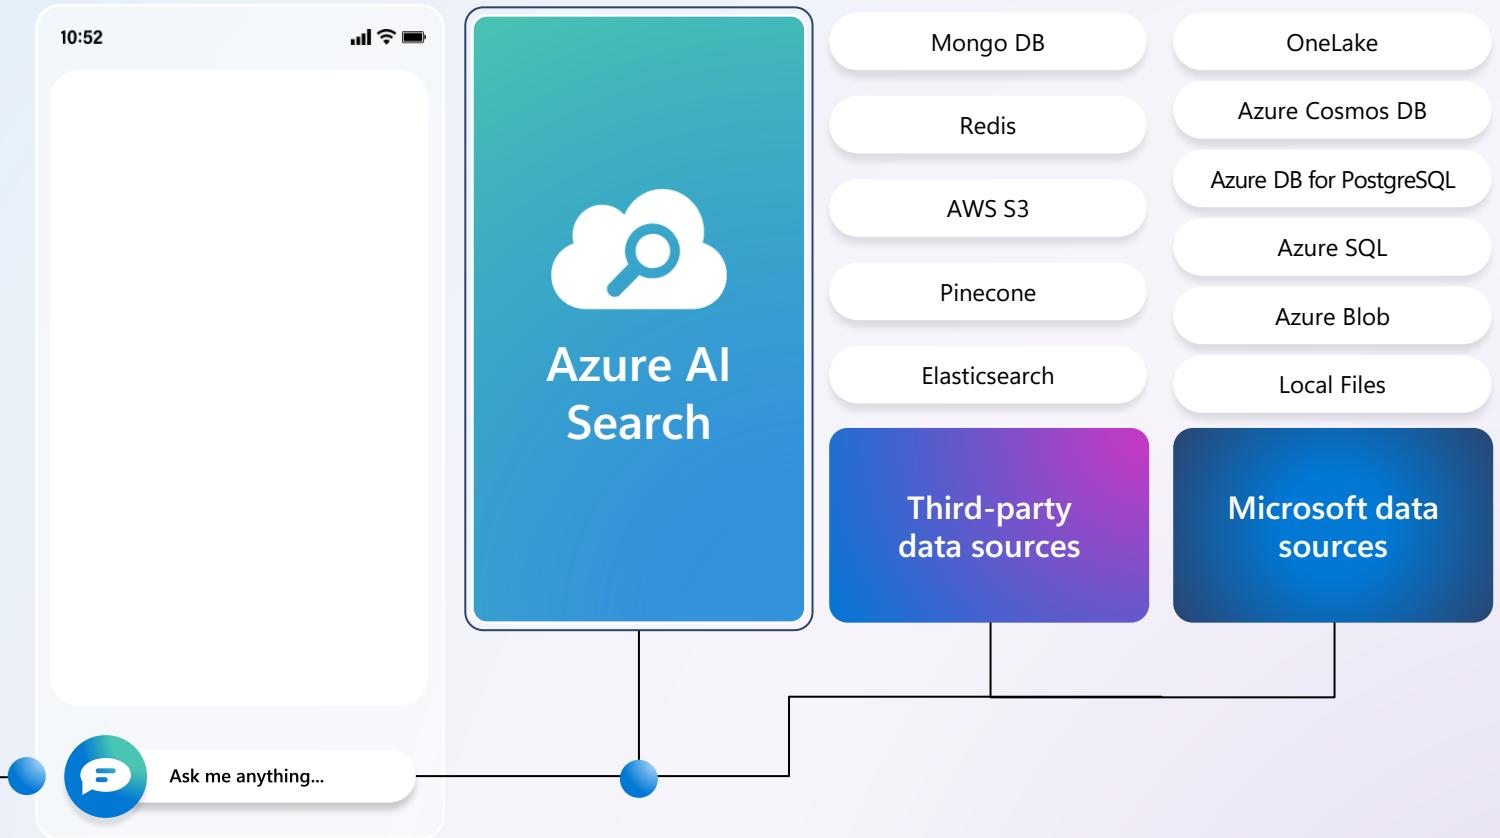
- Around human reading speeds- Try GPT4 Turbo, LLaMa3 70b

- Time response is irrelevant: Run the biggest best model you can on cloud deployments

What's your budget?

- Free : then go open-source
- Unknown, I'm testing an MVP: start with open-source, move to more performant models as needed
- Good price/performance ratio: OpenAI chatGPT 3.5. Its fast, cheap, and performs generally well.
- Money isn't a concern: Run the biggest best model you can on hosted hardware or pay for best in class API alternatives

Wherever your  
data resides, bring  
it to life through  
**Azure AI Search**



# Azure app platform

Deploy AI apps at scale with high performance

Developer-centric

Global scale

Code-to-cloud security

Any app, any language, any framework

Fully managed



Azure App Service



Azure Functions

Control and customization



Azure Container Apps



Azure Kubernetes Service

Integration services



Azure API Management



Azure Logic Apps

# Azure Databases power a new era of AI apps that run anywhere, at any scale.



## Intelligent

AI-powered management and performance at any scale



## Integrated

Modern, AI apps on an open, flexible, and common platform



## Trusted

A sustainable and reliable platform so your data is always available, safe and compliant

# Microsoft Fabric + Azure AI Foundry = boundless possibilities



**Microsoft Fabric**

Your unified  
solution for  
analytics

## Enhance data insights

Build enterprise chat applications to uncover insights using natural language from structured, unstructured, and real-time data



**Azure AI Foundry**

Your gateway to  
design, customize,  
and manage AI  
apps and agents

## Ground your AI on unstructured data

Ensure your AI is using the latest data in its interactions with your employees, partners, and customers

## Analyze customer interactions

Build speech analytics applications to enhance customer service, tailor support responses, and make data-driven decisions

## Customize machine learning models

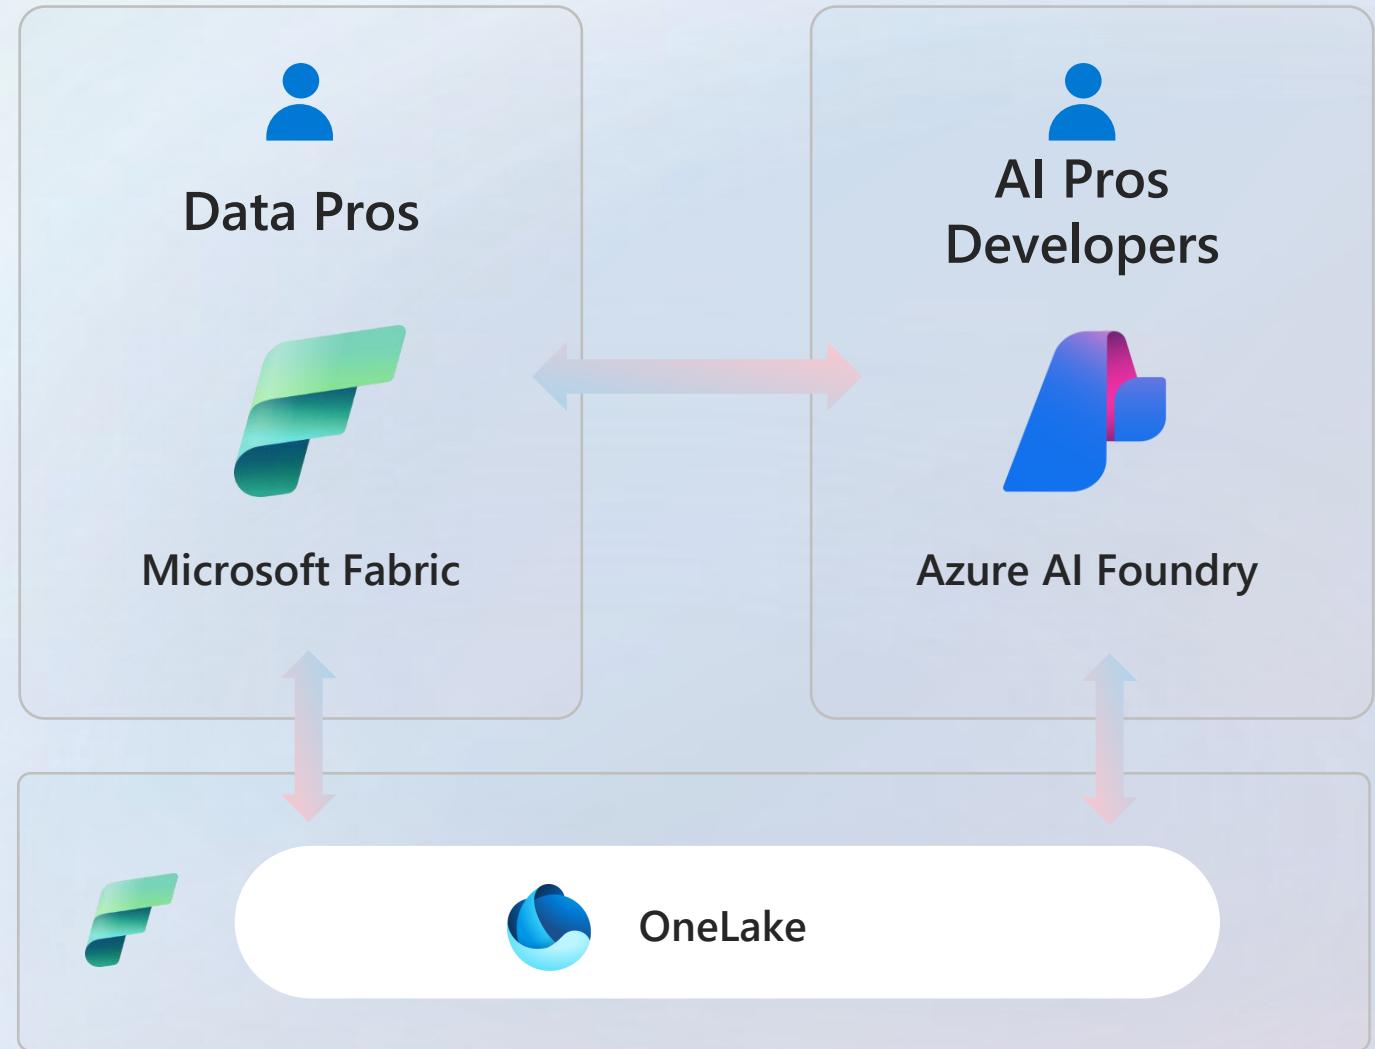
Train, deploy, and orchestrate machine learning models tailored to specific business needs – from predictive maintenance to customer sentiment analysis



OneLake serves as the connective tissue to build custom AI apps powered by your data

# Microsoft Fabric and Azure AI Foundry integration

- 1 Data and AI Pros access, use and process data across the AI lifecycle with OneLake as the connective tissue.
- 2 Data Pros can leverage Azure AI Search in Azure AI Foundry to index and search data stored in OneLake, enabling developers to build intelligent AI apps with up-to-date insights.
- 3 Developers can leverage Fabric data agent integration in Azure AI Foundry Agent Service as a knowledge source to build enterprise AI agents that can uncover actionable insights from structured data.



# Data agents in Microsoft Fabric

**Enhanced data interaction:** create conversational agents that understand the org's unique data context, making it easy for non-technical users to get reliable, data-driven answers.

**Faster time to insights:** significantly reduce the time it takes to generate, explore data, and discover trends by automating the process of transforming natural language questions into queries.

**Improved query capabilities:** train the data agent with example instructions and queries, connect up to 5 data sources, including Lakehouse endpoints and KQL databases, and get full visibility into each step taken by the agent.

Access within or outside of Fabric.

AI Application Developers



Azure AI Foundry  
Agent Service

Data Professionals

Publish/ Share data agent

Tools

Context and Configurations



OneLake



Microsoft Fabric



Lakehouse



Warehouse



Semantic  
Model



Eventhouse  
KQL DB

# Security for AI | Discovery Questions

- ✓ Identify existing practices and strategies to develop/evolve future Security for AI practices and improvements
- ✓ Leverage these discoveries to discern which security topics can be used for specific customer needs and sales team objectives

## Baseline Security

- **Where are you currently in your security journey to protect AI apps and agents? What are your biggest concerns about deploying AI agents?**
- This question will help understand what the customer knows about security solutions and helps capture excitement within the business strategy planning process. It sets a baseline for a conversation, and an overall foundation to start their journey.
- **Topics to consider:** Microsoft Security, Microsoft Zero Trust Solution

## Data leakage and inaccuracy

- **What data do you ground your AI agent on? Where does your sensitive data live? How do you prevent sensitive data from being exposed in AI prompts or responses? How are you protecting sensitive data used to ground your AI agent?**
- This helps identify whether Microsoft Purview is being used to classify, label, and prevent data loss. It also ensures quality data is being used to ground the AI app and any risky behavior is detected. It helps figure out whether AI usage insights are being flagged.
- **Topics to consider:** Microsoft Purview Data Security Posture Management for AI (DSPM for AI), Microsoft Purview SDK (label inheritance and real-time classification), Microsoft Purview Data Security (Information Protection, Data Loss Prevention, Insider Risk Management), Microsoft Purview Data Governance

## Emerging AI threats and vulnerabilities

- **Are you detecting and mitigating risks like prompt injection, model abuse, or jailbreak attempts in real time? Are you worried about adversarial attacks?**
- This helps assess whether the customer is using Defender for Cloud or Content Safety Prompt Shields protect against threats on models and vulnerabilities across the AI app stack. This also helps understand if the customer is using red teaming for AI systems.
- **Topics to consider:** Microsoft Defender for Cloud, Azure AI Content Safety Prompt Shields

## Agent sprawl and risks

- **How are you managing identity and access for your AI agents and applications? Do you have a centralized inventory of all AI agents deployed across your environments?**
- This helps uncover whether the customer is using Microsoft Entra to assign and govern non-human identities, and whether they have visibility into agent activity and access patterns. This helps assess the risk of agent sprawl and unmanaged access.
- **Topics to consider:** Microsoft Entra Agent ID

## Non-compliance

- **How are you preparing for AI-specific regulations like the EU AI Act or ISO 42001? What regulation is your organization subject to? How are you auditing your prompts and responses?**
- This helps assess whether the customer is using Microsoft Purview to generate compliance reports and audit trails as well as ensuring compliance with prompts and responses.
- **Topics to consider:** Microsoft Purview Compliance Manager and assessments, Audit, Data Lifecycle Management, eDiscovery, and Communication Compliance

# CISO Barriers | Objection Handling

- ✓ Identify potential barriers that a CISO might raise
- ✓ Understand how to sell and handle the objections using the respective products

## Risk of agent hallucination

- **Signals worry that agents might produce misleading or dangerous outputs**
- Prompt Shields and task adherence guardrails ensure grounded responses.

## Prompt injection and jailbreak attacks

- **Signals concern about malicious users manipulating AI**
- Prompt Shields protects against direct (user input) and indirect (tool output or data sources) prompt injection attacks and Defender helps detect, block, and mitigate the attacks. Real time threat detection and monitoring with Defender helps surface alerts and proves security visibility. These toolsets help prevent unintended or malicious prompt manipulation.

## Sensitive data leakage

- **Signals concern about agent inadvertently releasing sensitive information**
- The Purview SDK and Data Security Posture Management for AI help enable data protection policies, detect sensitive information, and surface recommendations to mitigate any sensitive data leakage.

## Agent sprawl

- **Signals worry about many agent instances**
- Entra Agent ID gives each agent an identity, enabling role-based access control and governance of agents.

## Adversarial Attacks

- **Concern about any attacks to the agent**
- Red Teaming simulates real world attacks. Azure AI Foundry also natively integrates security posture checks, guardrails, and evaluations allowing developer and security teams to work together closely.

# Secure and govern AI with Microsoft

Start secure



Stay secure

## Security and governance capabilities

### Protect secret and code

GitHub Advanced Security

### Build data security and safety guardrails

Azure AI Content Safety and Microsoft Purview

### Pentest and evaluate AI

AI Red Teaming Agent and evaluation

### Manage posture and protect against threats

Microsoft Defender

### Secure and govern data and access

Microsoft Purview and Entra

## Foundational commitments

### Secure by design, secure by default, secure operations

Secure Future Initiative

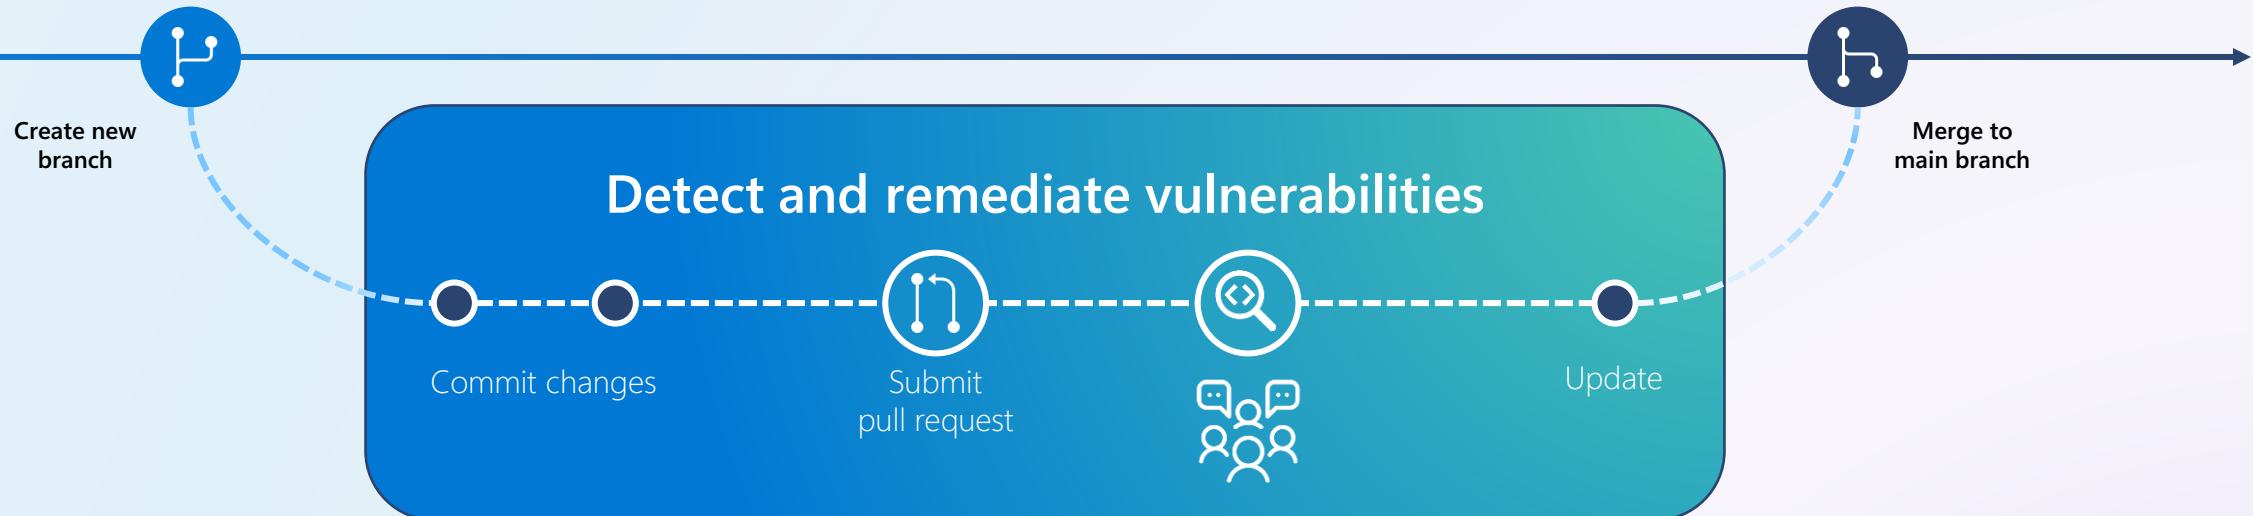
### Data is private at work, at home, and on the go

Privacy principles

### Fairness, reliability and safety, privacy and security, inclusiveness, transparency, accountability

AI principles

# Shift security left with GitHub Advanced Security



Powered by Copilot



## GitHub Secret Protection

- Secret scanning
- Copilot Secret Scanning
- Push protection
- Secret risk assessment

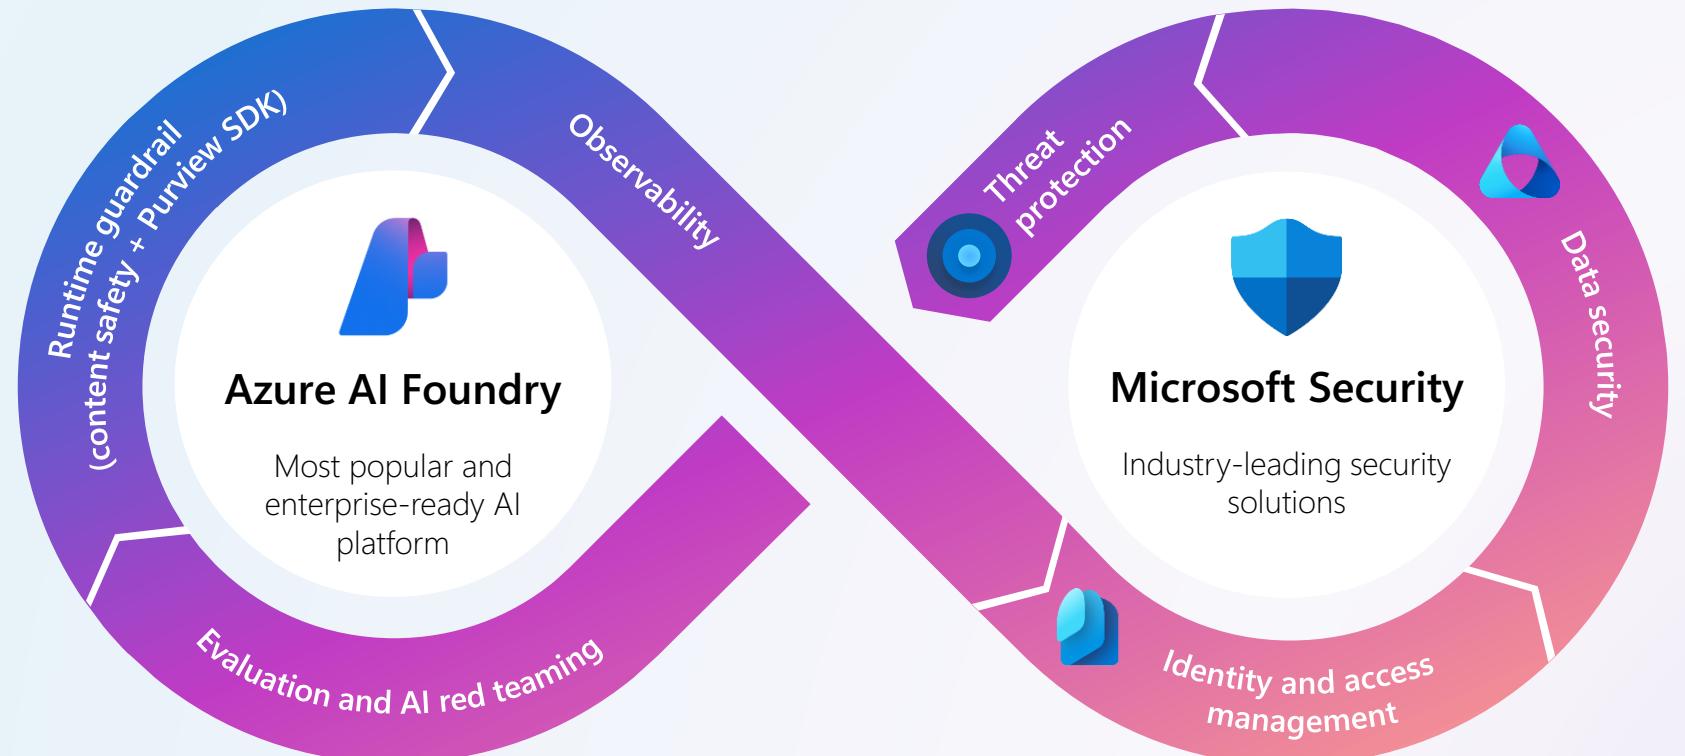
Powered by Copilot



## GitHub Code Security

- Code scanning
- Code quality
- Copilot Autofix and Security Campaigns
- Dependency graph and review
- Dependabot

# Securing and governing AI from code to cloud is a team sport



Shared responsibility to start secure and stay secure



# Integrations to secure AI apps and agents



## Microsoft Entra Agent ID

Manage agent identities  
and lifecycle

Secure and govern agents'  
access to resources

Deploy out of the box with Copilot Studio  
and Foundry



## Purview integration with Azure AI Foundry

Discover data security  
and compliance risks

Protect against data leaks  
and insider risks

Enable governance for AI and meet  
regulatory needs



## Defender and Azure AI Foundry

Get alerts and recommendations from  
Defender directly in Foundry

Streamline security checks

Proactively identify and address security  
vulnerabilities during development

# Azure AI Content Safety



Harm

**Content filters** for text and images in inputs/outputs:

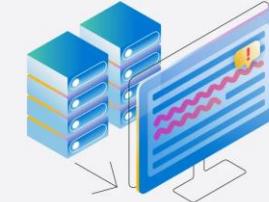
- Violence
- Hate and fairness
- Sexual
- Self-harm
- Custom categories



Security

**Prompt shields** for inputs:

- Direct prompt injection attacks
- Indirect prompt injection attacks



Quality

**Detection models** for inputs or outputs:

- Groundedness
- Protected text
- Protected images
- Protected code
- Custom blocklists

# Azure AI Foundry Observability

Continuously improve and monitor AI agents

## Balance impact, quality, safety, and cost

Total requests  
**97.4K**

Total token cost ⓘ  
**\$3,203**

### Operational metrics

Total tokens  
**6,598**  
+5% increased over last 7 days

Latency  
**6.2s**  
+5% increased over last 7 days

### Evaluation metrics

Groundedness  
**82%**  
+5% increased over last 7 days

Coherence  
**95%**  
+5% increased over last 7 days

Jailbreak  
**21%**  
-5% decreased over last 7 days

Hallucination  
**17%**  
-5% decreased over last 7 days

## Develop with confidence

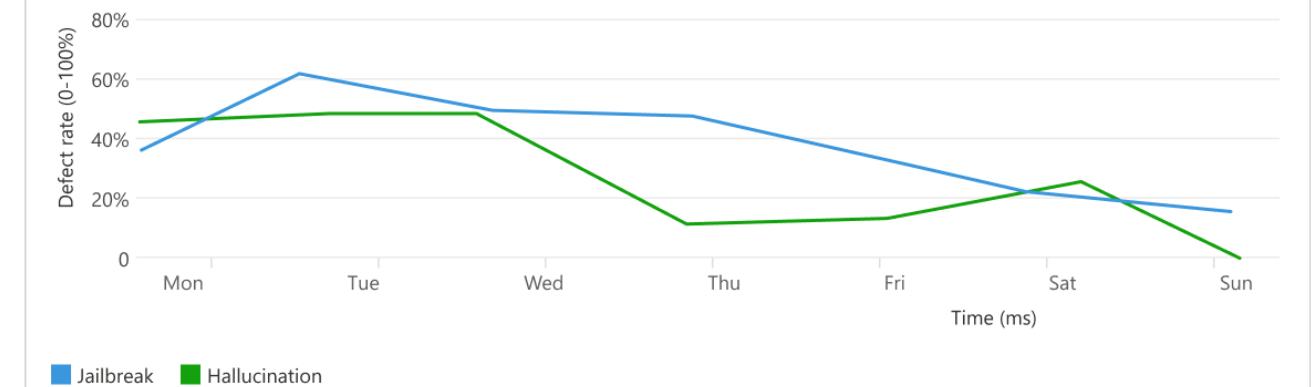
All checks have passed  
3 successful checks

✓ ⚡ [AI Evaluation / Test Variants with AI Evaluations \(push\)](#)

Evaluation score	my-agent-v1	my-agent-v2
Groundedness	Baseline 5.00	Degraded 1.40 (-3.60)
Relevance	Baseline 4.90	Inconclusive 4.60 (-0.30)

## Operate and monitor in production

≡ Daily safety defect rate





# Industry compliance

- ✓ CSA STAR Certification
- ✓ FedRamp
- ✓ GDPR
- ✓ GSMA SAS-SM
- ✓ GxP
- ✓ HDS
- ✓ HIPAA BAA
- ✓ HITRUST
- ✓ ISO 9001, 20000, 22301, 27001, 27017, 27018, 27701
- ✓ PCI 3DS, DSS
- ✓ Australia IRAP
- ✓ Spain ENS High
- ✓ Singapore OSPAR
- ✓ UK Cyber Essentials Plus
- SOC – coming soon!

# Why innovate with AI apps and agents on Azure

Microsoft Azure empowers organizations to **build, customize, scale, and secure AI apps and agents** with unmatched speed, flexibility, and trust

## Accelerate innovation

Use GitHub Copilot and Azure DevOps to enhance developer velocity and automate routine coding, testing, and deployment

## Stay flexible and future-ready

Design, customize, and manage AI apps and agents seamlessly with 11,000+ models and a comprehensive multi-agent toolchain in Azure AI Foundry

## Scale with confidence

Deploy apps globally using Azure's cloud-native services, reusable templates, and robust infrastructure (99.99% uptime)

## Ensure trust and control

Rely on end-to-end security, governance, and safety controls built on foundational platform commitments to secure and govern AI throughout the AI development lifecycle

# Key Action: Download Sales Assets from aka.ms/AIForPartners/CoSell/FY26

FY26 - Innovate with AI Apps and Agents Partner Package

Last Modified 2025-09-17

Use the content in this package in customer conversations for success through the play.



This campaign's content

Showing 1-12 of 15 assets

Asset	Description	Date	Category
Agentic DevOps L200 deck - For Partners	L200 deck for Agentic DevOps.	2025-09-17	Intelligent Cloud
AI Executive (Halo) Conversation - Becoming Frontier	AI Exec Pitch Deck.	2025-09-17	Intelligent Cloud
Azure AI Foundry L150 Pitch Deck	Pitch Deck for Azure AI Foundry.	2025-09-17	Intelligent Cloud
Azure AI Foundry L300 Pitch Deck	Deep dive deck for Azure AI Foundry.	2025-09-17	Intelligent Cloud
Azure AI Agent Service L200 Pitch Deck	Deep Dive Customer Presentation for Azure AI Agent Service.	2025-07-28	Intelligent Cloud
Azure AI-optimized Databases L150 pitch deck for databases for AI use-cases.	Github Repo - Build your first agent.	2025-09-17	Intelligent Cloud
GitHub Repo - Build your first agent	Github Repo - Build your first agent.	2025-07-28	Intelligent Cloud
GitHub Repo - AI agents for beginners	Lessons to start building AI agents.	2025-07-28	Intelligent Cloud
Innovate with AI Apps and Agents L200 Deck - For partners	Deep Dive Customer Presentation for Azure AI Agent Service.	2025-07-28	Intelligent Cloud
Innovate with AI Apps and Agents L200 Deck - For partners	Deep Dive Customer Presentation for Azure AI Agent Service.	2025-07-01	Intelligent Cloud
Gold Standard Solution Accelerators	Microsoft Solution Accelerators.	2025-07-01	Intelligent Cloud
GitHub Copilot Enterprise Demo	GitHub Copilot Enterprise Feature Walkthrough Demo.	2025-07-01	Intelligent Cloud
Innovate with AI Apps and Agents L200 Deck - For partners	Deep Dive Customer Presentation for Azure AI Agent Service.	2025-07-01	Intelligent Cloud

Showing 1-12 of 15 assets

1 2 Next >

Select All



**Connect with  
developers with  
GitHub and drive  
them to Azure**



# Opening Agentic DevOps Conversations

**Guidelines for the pitch:** Focus on developer productivity. Emphasize AI integration across the full platform, especially integrations with Azure AI Foundry.

Persona	How to start the conversation
Decision Makers	Focus on faster time to market, improved code quality, improve developer experience, giving back developers time to upskill with rapid changes in the AI space, and better use of developer skill sets.
Practitioners	Aim to focus on challenges pertaining to repetitive tasks, such as test cases, debugging inefficiencies, context switching, etc. Focus on highlighting how GitHub Copilot, and its integrations with Azure AI Foundry can mitigate these challenges.

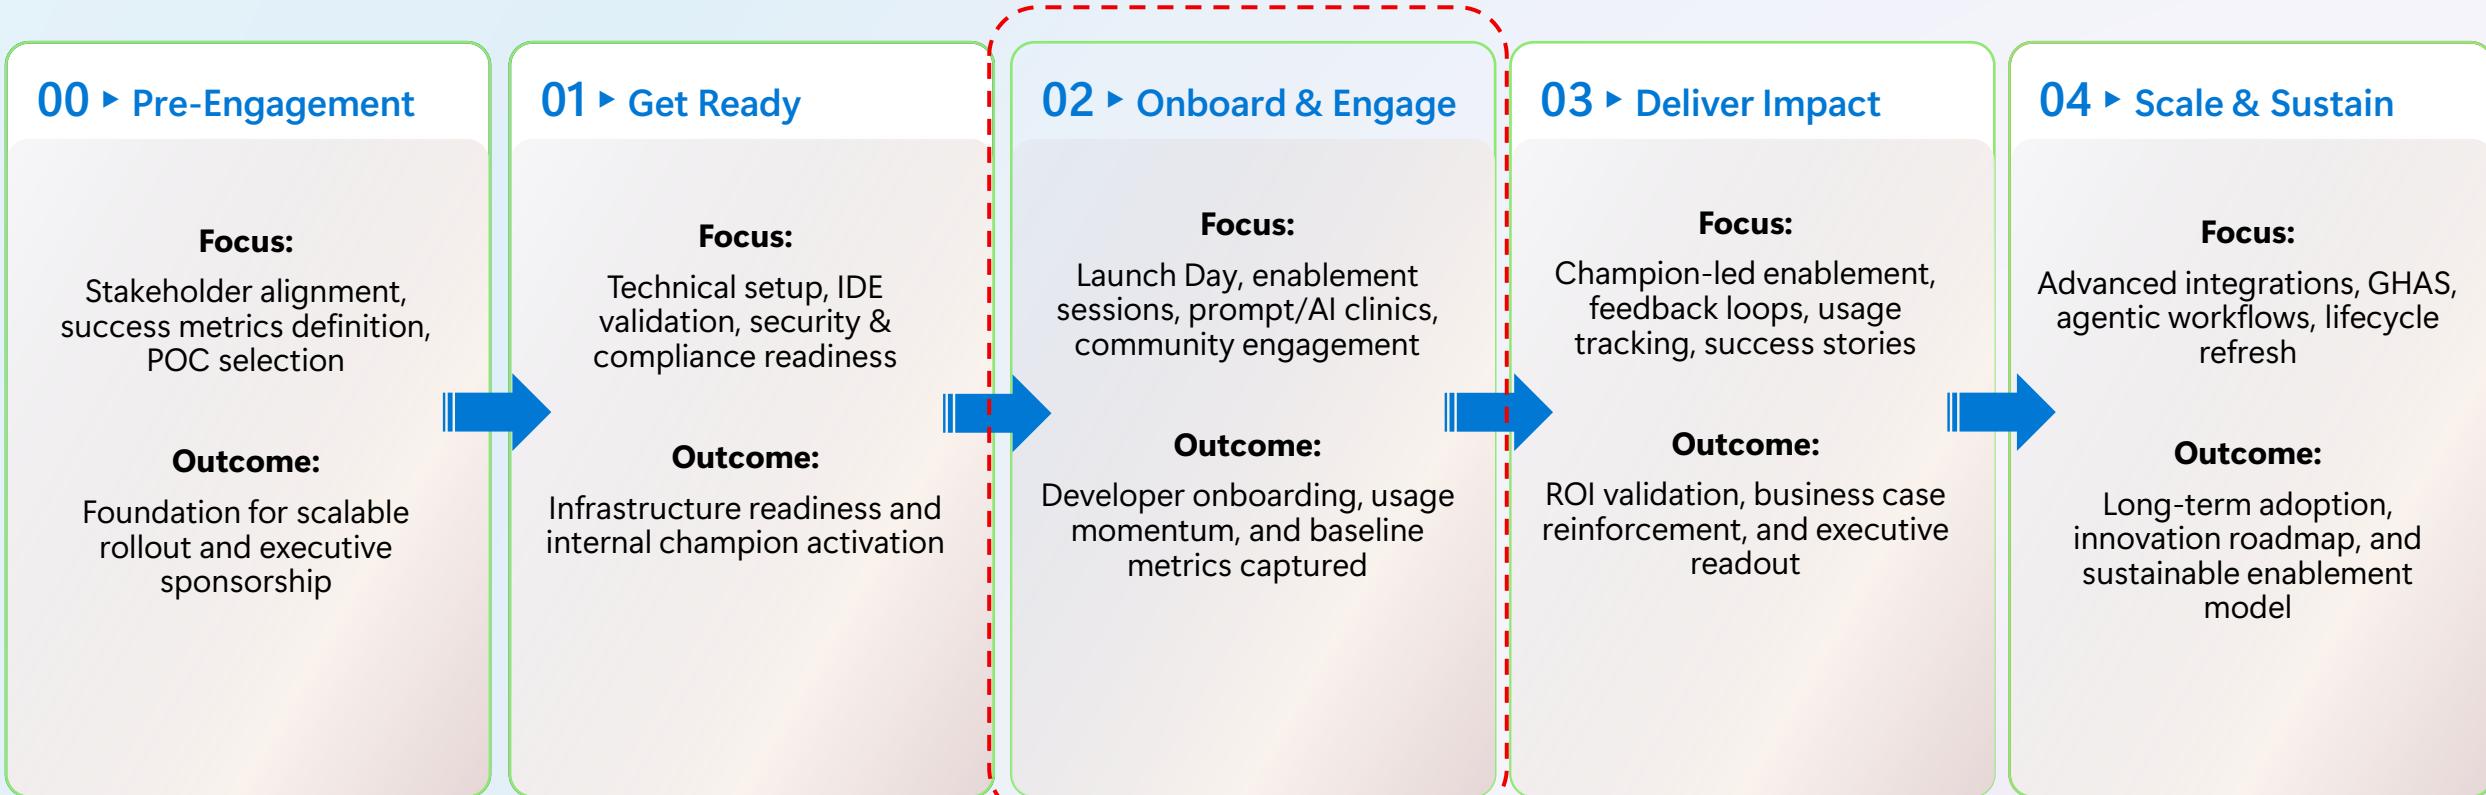
## Discovery questions:

- Given the increasing pace of AI app and agent development, what are your plans to get your developers to do more, faster?
- Have you explored autonomous agents or GitHub Copilot beyond code completion—for example, in planning, testing, or documentation?
- Do you have governance structures in place for AI or agentic workloads?
- What is your current process to fix security vulnerabilities? Would it help to have an AI agent assist with that?
- Do you need tools to help you improve your test coverage?
- When trying to build an AI app or agent, how often are you switching between tools?
- How are you preparing your teams for AI-native development? Are there upskilling or reskilling initiatives underway?

## Key resources:

[Agentic DevOps L100 Pitch Deck](#) | [Agentic DevOps L200 Pitch Deck](#)

# GitHub Copilot Adoption – Strategic Journey



Scalable as per the revenue potential opportunity

# GitHub Copilot Adoption – Milestones

## 00 ▶ Pre-Engagement

- Identify key customer stakeholders
- Define success metrics and segment developer cohorts.
- Scope rollout phases and align with business goals.
- Form GitHub Copilot Champions Team and schedule kickoff.
- Select a POC project with measurable outcomes and strategic relevance

## 01 ▶ Get Ready

- Validate technical readiness (IDE, proxy, SSO, metrics API).
- Approve security and compliance protocols.
- Build enablement strategy (Comms, Launch day, Enablement sessions).
- Form internal Copilot Guild and kick off POC.
- Share early POC insights in internal comms.

## 02 ▶ Onboard & Engage

- Deploy Copilot and activate usage dashboards.
- Deliver Launch Event & Partner-led enablement sessions using live POC examples.
- Train Champions and activate internal Copilot Guild.
- Launch AI Clinics and mentoring formats.
- Define CoE roles and track adoption stages.

## 03 ▶ Deliver Impact

- Run surveys and usage reports to validate impact.
- Finalize POC and showcase results.
- Champion-led Enablement sessions for continued adoption
- Deliver executive readout and POE.
- Address adoption gaps with refreshers and contests.

## 04 ▶ Scale & Sustain

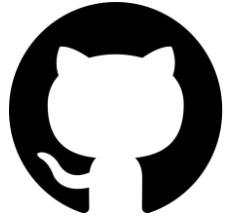
- Expand Copilot into advanced features, API integrations, agentic workflows and new teams.
- Complete platform and GHAS integrations.
- Launch partner-led innovation and optional modules.
- Plan lifecycle refresh and compete strategy.

OPTIONAL

Scalable as per the revenue potential opportunity

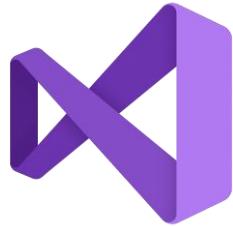
# Connecting with Developers – Product Highlights





## GitHub

**150M+** accounts



## Visual Studio

**50+ million**  
monthly engaged users



## Copilot Studio

**90%** Fortune 500 companies

World's **most loved** developer tools

# Agentic DevOps

AI-powered agents operating as a member of your dev team—automating, optimizing, and accelerating every stage of the software lifecycle



## GitHub Copilot

### Agent Mode

AI-powered pair programmer, performing multi-step tasks at the developer's command

### Coding Agent

AI-powered peer programmer, automating routine tasks and jobs to be done across the software lifecycle

### App Modernization

GitHub Copilot extends AI-assisted capabilities to accelerate modernization and upgrades for .NET and Java apps

Public Preview



## SRE Agent

AI-powered operator, continuously monitoring and optimizing apps in production

Public Preview



## GitHub Models

Every developer can be an AI developer with full model accessibility and fine-tuning in GitHub to speed app development

Public Preview



# GitHub Copilot is built for the developer experience

Automate the mundane and empower developers to focus on what matters

## Experiences

Sense making (summarization, code understanding)

Writing code

Writing tests, documentation

Debugging code

Reviewing code

Fixing bugs and vulnerabilities

Code modernization, refactoring

Upgrading dependencies

Updating stakeholders

# Visual Studio family



## Most-loved AI-infused developer tools

## Built for flexibility

Any platform, any language, any device—from local development to cloud-scale deployments

 Integrated AI tooling with GitHub Copilot

AI-powered coding assistance that reasons through problems, coordinates next steps, applies changes, and iterates on errors

 Azure AI platform integration

Create cloud-first applications efficiently and securely using integrated tools and AI capabilities without leaving the IDE

File Edit Selection View Go ... 🔍 TODO app - Copy 🌐

Agent (Prompt) Builder

Recent + New Builder

what the time and we...

Model \* Settings

OpenAI GPT-4o (via GitHub)

Prompts \* Add Prompt

Click the button below to automatically generate a prompt with your own description, or enter your prompts manually.

Generate system prompt

System prompt ⓘ

You are an AI assistant designed to help users manage their task lists. Your capabilities include adding, updating, and deleting tasks. Additionally, you are responsible for recommending tasks by evaluating priority, due dates, current time, and weather conditions. Do all tasks through these 3 MCP servers: todo, weather, time. Save and read user tasks in "/Users/ronglu/aitk-example/TODO.json".

# Steps

1. \*\*Task Management Actions:\*\*

- For adding tasks: Accept user inputs such as task name, due date, priority, and any additional details required.
- For updating tasks: Prompt the user for the task's identifier and the fields they want to update (e.g., name, priority, due date).
- For deleting tasks: Validate the task to delete based on its identifier or unique parameters. Confirm action with the user if necessary.

2. \*\*Recommendation Logic:\*\*

"is\_completed": false  
}

Model Response

Based on the current context—weather and tasks on your list—here's a recommendation for your next activity:

json

{ "Recommendations": [ { "Current Date and Time": "2025-05-09T10:48:05-07:00", "Weather": "87°F, mostly sunny", "Task": "Prepare for //build talk", "Reason": "High priority and overdue since April 29, suitable for the current indoor conditions given the warm weather." } ] }

Would you like to proceed with this task or need suggestions on a different one?

Token used: 130

Run

View Code



# Copilot Studio

Build AI apps and agents  
that **extend Microsoft 365 Copilot** or  
**operate standalone**

The screenshot shows the Copilot Studio web interface. At the top, there's a search bar with placeholder text "Describe your agent to create it" and a button "Create". Below the search bar, there's a section titled "Recent" with a table showing two items:

Name	Type	Modified by	Last published	Owner
Microsoft 365 Copilot	Microsoft	Mona Kane	2 weeks ago	Mona Kane
Contoso Bank	Agent	Pedro Armijo	2 months ago	Mona Kane

Below the recent section is a "Start with a template" section featuring six cards:

- IT Helpdesk: Empowers employees to resolve issues and effortlessly create/view support tickets.
- Safe Travels: Provides employees the travel assistance in terms of travel inquiries and health and safety guidelines.
- Store Operations: Improve the efficiency of retail frontline workers by enabling easy access to store procedures and policies.
- Sustainability Insights: Enables users to easily get insights and data about a company's sustainability goals and progress.
- Team Navigator: Assists employees in finding colleagues and their hierarchy within the organization with Microsoft Tea...
- Voice: An agent with voice capabilities.

At the bottom, there's a "Learning resources" section with several cards:

- Quick start: Create and deploy an agent
- Documentation
- Responsible AI FAQs
- Quick start: Use Generative AI in an agent
- Support community
- Try autonomous agents

Meet users where they already are

Access everything in one place

Automate workflows

Integrate with external apps

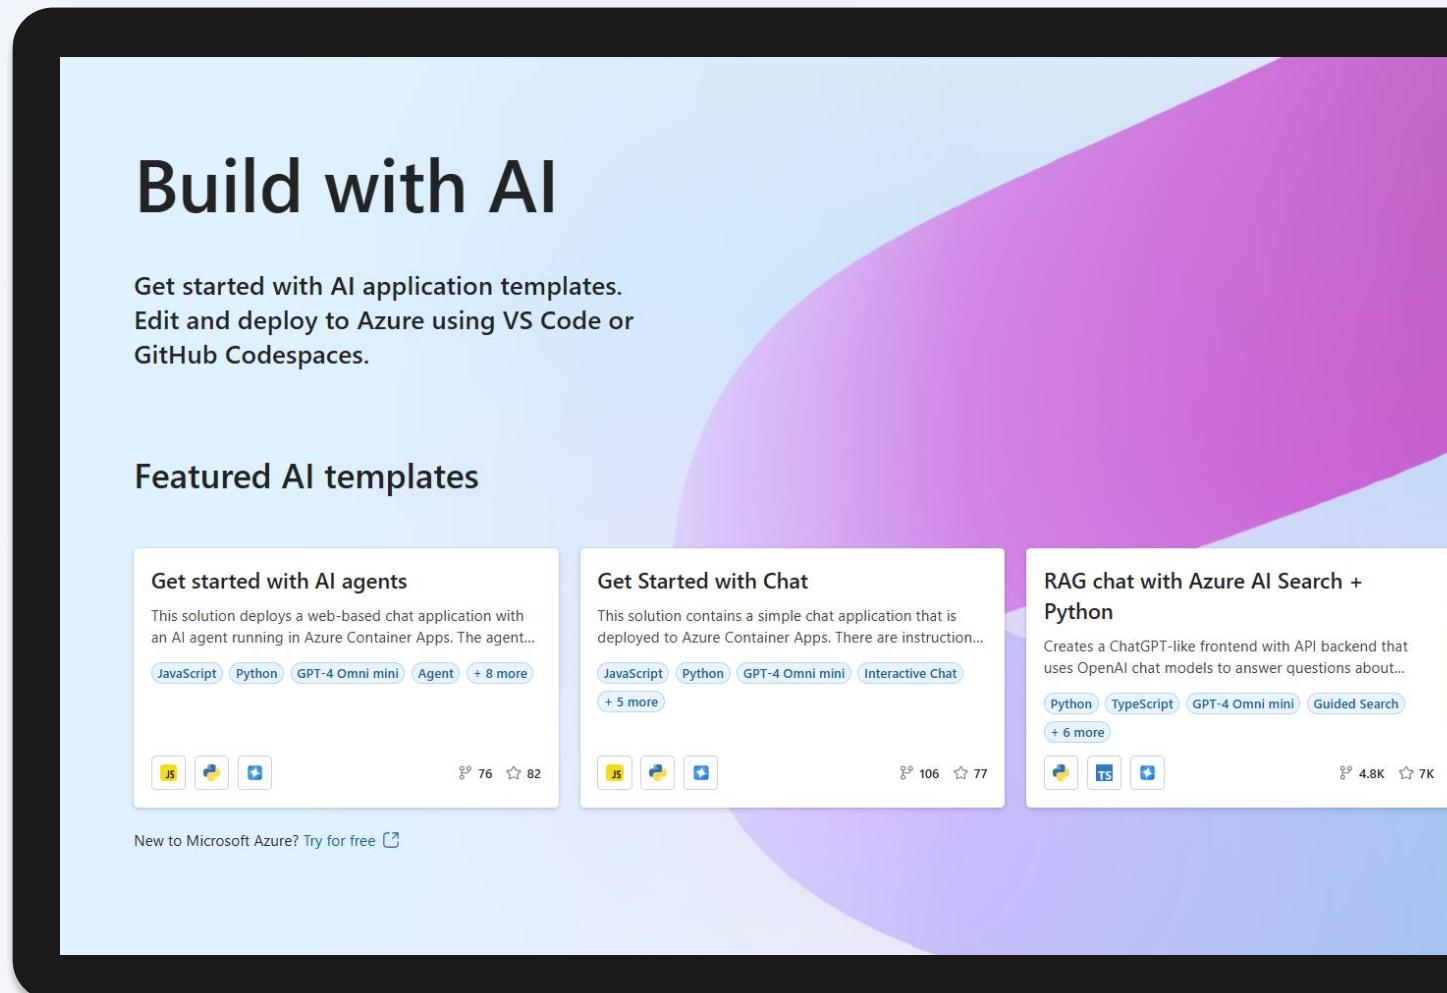
Connect to data in Microsoft 365

# AI app template library

Empowering developers to easily build AI apps on Azure

- Curated repos for common AI app patterns
- Codified best practices for evaluating, deploying, and monitoring production-quality apps
- Repos that include both Azure-native stacks and stacks based on popular community projects
- Includes customizable infrastructure as code assets and CI/CD pipelines
- Deploy each sample app to Azure using Visual Studio Code and GitHub Actions

<https://aka.ms/ai-apps>



# Key Action: Download Sales Assets from aka.ms/AIForPartners/CoSell/FY26

FY26 - Innovate with AI Apps and Agents Partner Package

Last Modified 2025-09-17

Use the content in this package in customer conversations for success through the play.



This campaign's content

Showing 1-12 of 15 assets

Asset	Description	Date	Category
Agentic DevOps L200 deck - For Partners	L200 deck for Agentic DevOps.	2025-09-17	Intelligent Cloud
AI Executive (Halo) Conversation - Becoming Frontier	AI Exec Pitch Deck.	2025-09-17	Intelligent Cloud
Azure AI Foundry L150 Pitch Deck	Pitch Deck for Azure AI Foundry.	2025-09-17	Intelligent Cloud
Azure AI Foundry L300 Pitch Deck	Deep dive deck for Azure AI Foundry.	2025-09-17	Intelligent Cloud
Azure AI Agent Service L200 Pitch Deck	Deep Dive Customer Presentation for Azure AI Agent Service.	2025-07-28	Intelligent Cloud
Azure AI-optimized Databases L150 pitch deck for databases for AI use-cases.	Github Repo - Build your first agent.	2025-09-17	Intelligent Cloud
GitHub Repo - Build your first agent	Github Repo - Build your first agent.	2025-07-28	Intelligent Cloud
GitHub Repo - AI agents for beginners	Lessons to start building AI agents.	2025-07-28	Intelligent Cloud
Innovate with AI Apps and Agents L200 Deck - For partners	Deep Dive Customer Presentation for Azure AI Agent Service.	2025-07-28	Intelligent Cloud
Gold Standard Solution Accelerators	Microsoft Solution Accelerators.	2025-07-01	Intelligent Cloud
GitHub Copilot Enterprise Demo	GitHub Copilot Enterprise Feature Walkthrough Demo.	2025-07-01	Intelligent Cloud
Innovate with AI Apps and Agents L200 Deck	L200 deck for Innovate with AI Apps and Agents.	2025-07-01	Intelligent Cloud

Showing 1-12 of 15 assets

1 2 Next >

Select All

# Fueling cloud and AI transformation

One unified offering bringing together....

## Azure Innovate

Infuse AI into solutions and advance analytics capabilities

## Azure Migrate and Modernize

Securely and efficiently move existing workloads to Azure

Announcing

# Azure Accelerate

Backed by foundational resources and best practices...

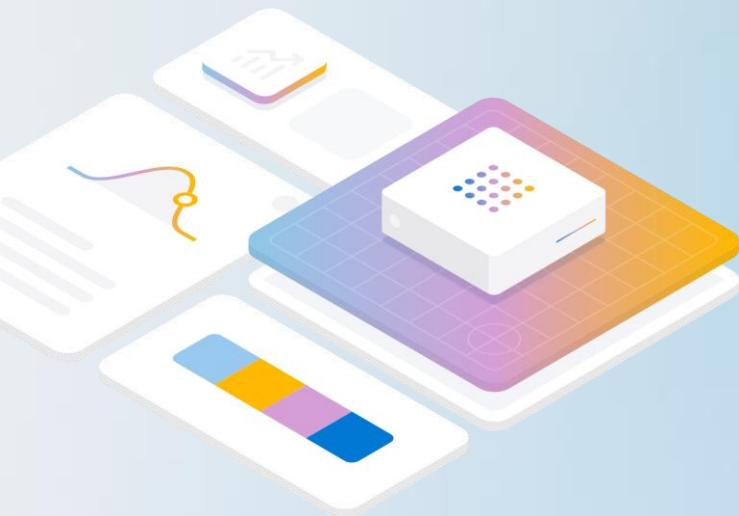
## Azure Essentials

Elevate reliability, security, and ongoing performance of cloud and AI investments with extensive guidance, resources, and tooling

[aka.ms/azureploffering](http://aka.ms/azureploffering)

# Azure Accelerate

## What's new in FY26



### One unified offering

- One **brand** for all customer needs across migration, modernization and innovation
- A **simplified offer construct** with straightforward funding across solution plays and project sizes
- Clear **project size criteria** between partner-nominated and Microsoft field-nominated projects

### Fast approvals via automation

- A streamlined partner nomination experience with the **latest offers reflected in MCI**

### Richer partner investments

- Increased **partner funding** for pre-sales and select strategic workloads (VMware migrations and SAP)
- Integration of **Cloud Accelerate Factory** joint delivery at no additional cost
- Solutions **assessment approach** enabled by partner access to 1<sup>st</sup> party and 3<sup>rd</sup> party tools like Dr. Migrate

### Increased sales velocity

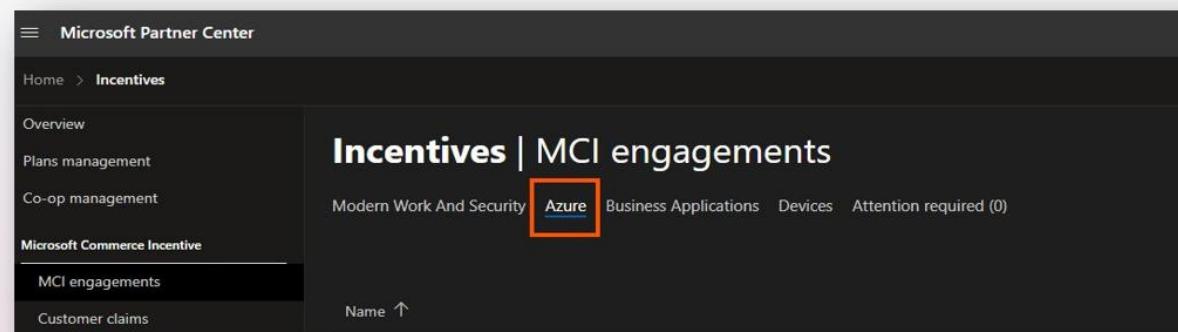
- Packaged pre-sales offer includes **assessment and proof of value** in a single engagement
- Go-to-market support to build pipeline and win more deals for **enterprise and SMB customers**

# Microsoft Commerce Incentives (MCI)

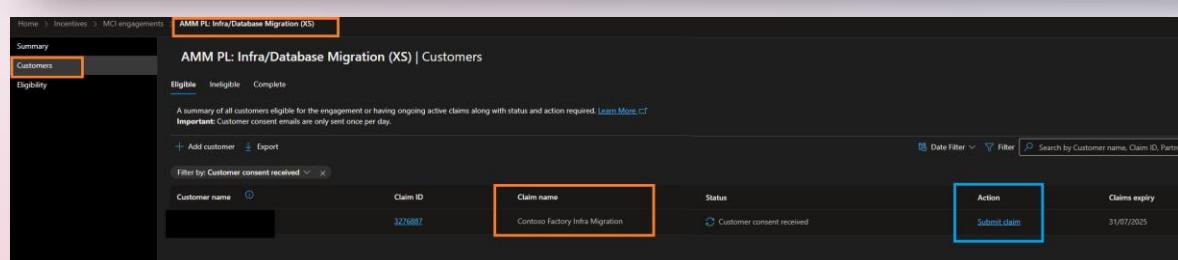
Partner nominated front door to Azure Accelerate



## Claiming the customer in MCI



The screenshot shows the Microsoft Partner Center interface. On the left, there's a sidebar with links like Home, Overview, Plans management, Co-op management, Microsoft Commerce Incentive (which is expanded), MCI engagements (highlighted with a black box), and Customer claims. The main content area is titled "Incentives | MCI engagements". Below the title, there are tabs for Modern Work And Security, Azure (highlighted with an orange box), Business Applications, Devices, and Attention required (0). A search bar at the bottom is labeled "Name ↑".



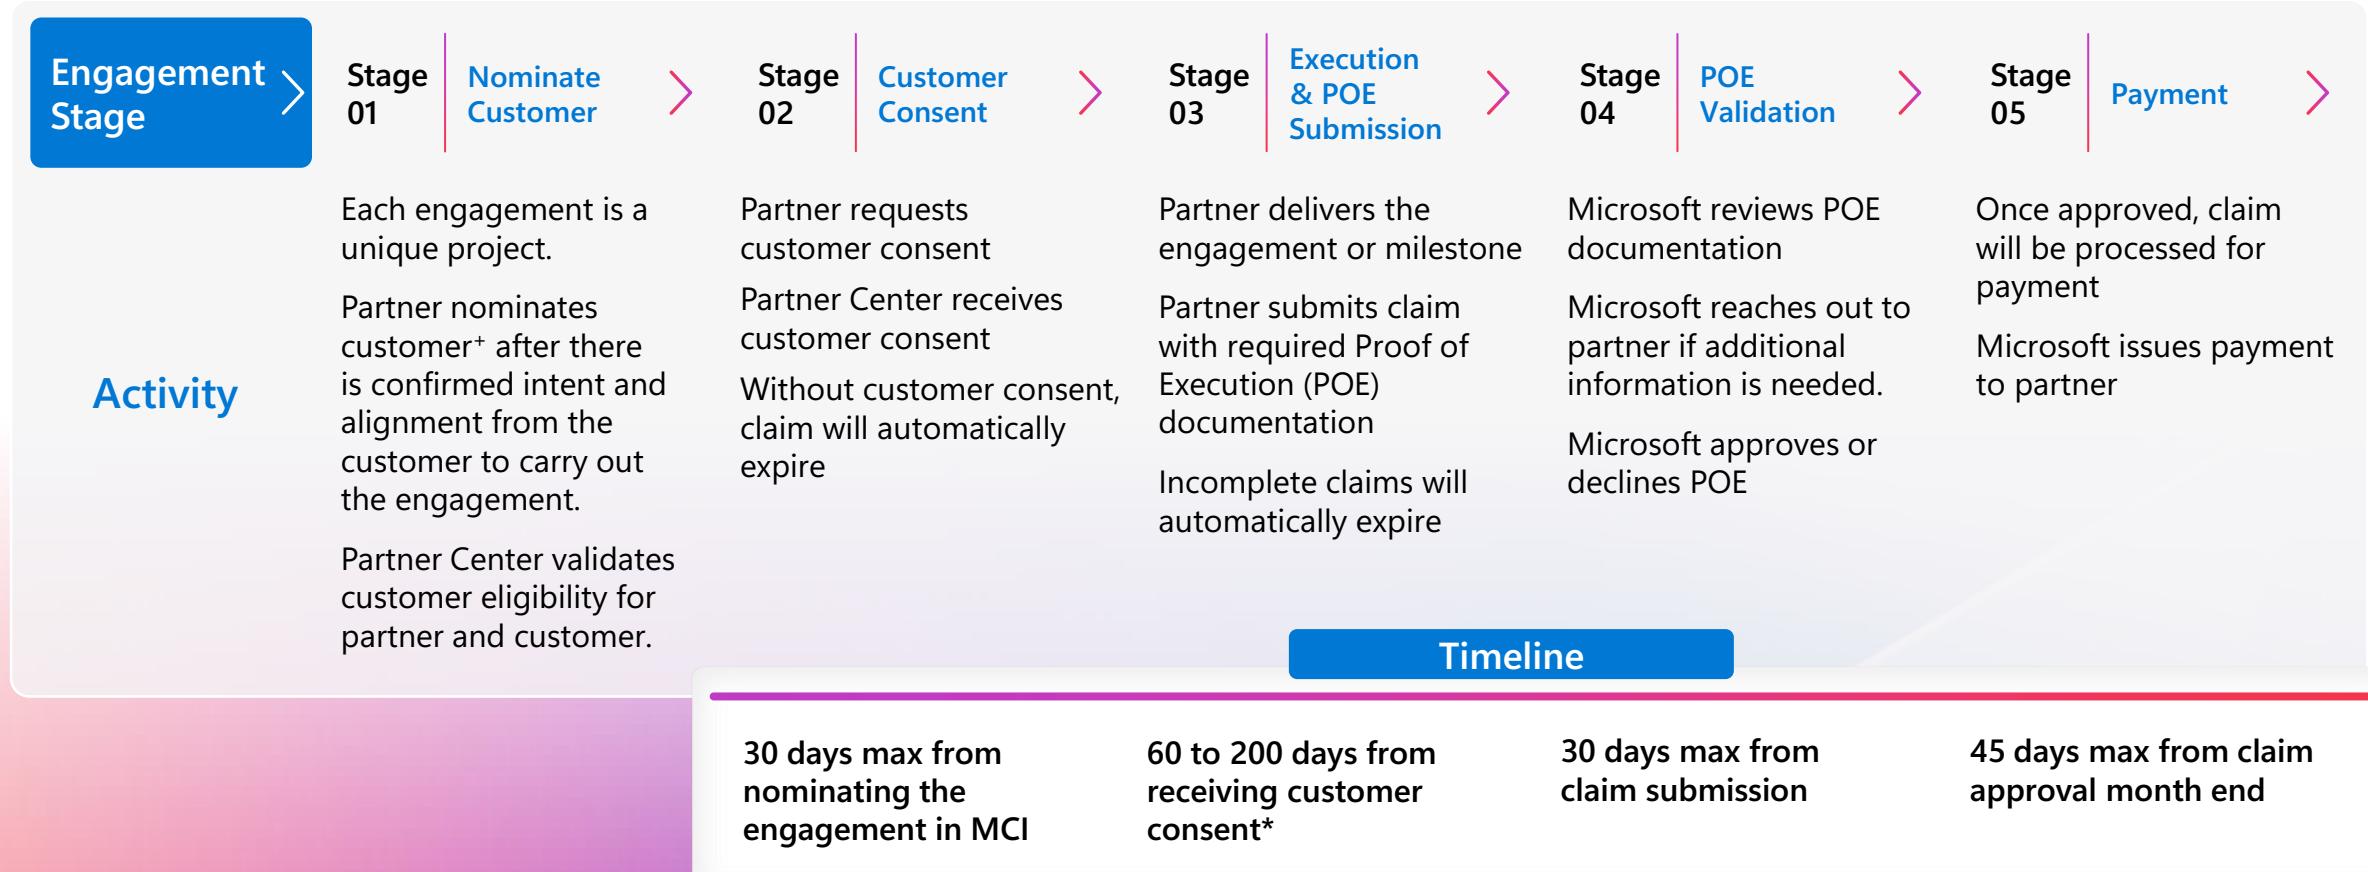
This screenshot shows a detailed view of a specific engagement. The top navigation bar includes "Home > Incentives > MCI engagements > AMM PL: Infra/Database Migration (XS)". The left sidebar has tabs for Summary, Customers (highlighted with an orange box), and Eligibility. The main content area displays a table with columns: Customer name, Claim ID, Claim name, Status, Action, and Claims expiry. One row is visible, showing "Contoso Factory Infra Migration" under Claim name and "Customer consent received" under Status.

Navigate to the incentive's portal directly  
<https://aka.ms/AzurePLOfferingsNominate>  
& click on the Azure tab

To close out and submit the claim, follow the same steps as before to navigate to the engagement & search for the claim by Customer Name or the Claim Name you assigned to the project. Then click on the "Submit Claim" button.

# Preparing and nominating your customer

## Partner nominated pathway



+ Nominate by clicking "Add Customer" in Partner Center > Incentives > MCI Engagements. See <https://aka.ms/MCIPartnerActivitiesClaimsGuide> for details.

\*The exact timeline for this stage will depend on the size of the engagement.

Please refer to the individual one-pagers of the respective engagement type for more details

# Partner expectations and requirements



## Ensure Customer intent before submitting request

To truly take advantage of the new process at scale, Partners should ensure that their projects or customers are fully aligned and committed **before** submitting the request in the MCI Portal



## One funding instance request per customer project

Partners may not use funding for the same project scope across multiple sources. This also includes leveraging Cloud Accelerate Factory delivery services.

# Azure governance and partner status reporting



## Governance requirements

Partners are required to provide accurate Azure Subscription ID information (where applicable): We will be monitoring the Azure Subscription IDs very closely and tracking consumption.

Partners who do not meet a [minimum success criteria outlined here](#) might be 'paused' from nominating future engagements.

[Learn more here about Governance](#)



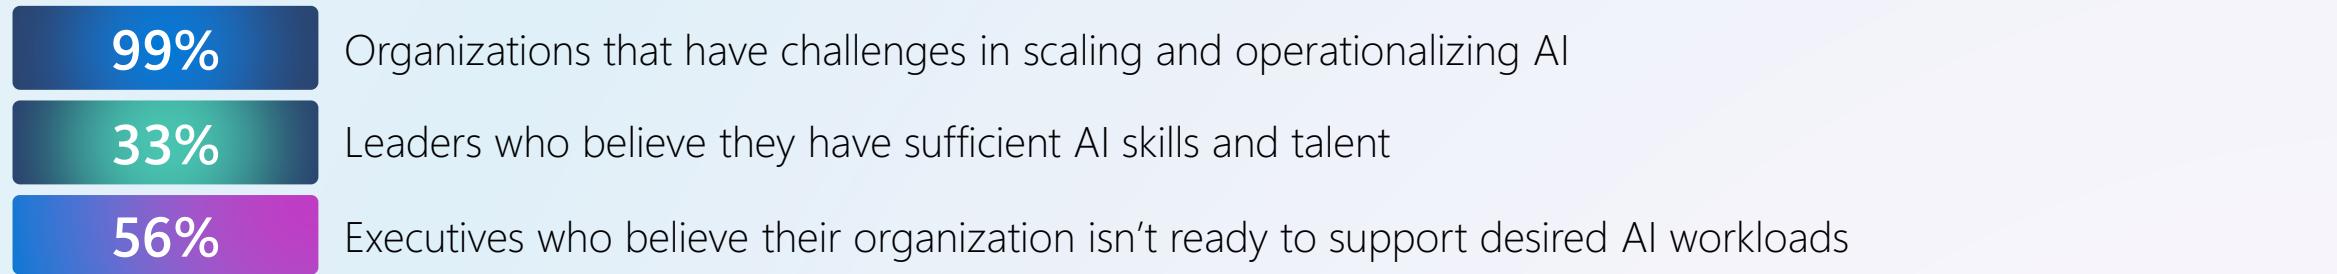
## Partner Status Report

Partners can now access regular status report on their participation in Azure Accelerate globally for their organization delivered straight to their inboxes twice a month!

[Sign up now:](#)  
<https://aka.ms/AzurePLOfferingsReportSignup>

\*Refer to the [Microsoft Commerce Incentives Guide](#) for complete terms and conditions.

# Best practices guidance and tools with Azure Essentials



- Financial best practices
- Pricing
- Skilling
- Workload migration and data collocation
- AI Center of Excellence
- Platform landing zone
- Trustworthy AI: security, privacy, and transparency
- Architectural guidance for AI workloads
- GenAIOps
- Data governance
- Monitoring and security
- Management and optimization
- Architectural assessments and remediations

## Key Resource Links:

[aka.ms/azureploffering](https://aka.ms/azureploffering)

[aka.ms/AzureEssentialsPartners](https://aka.ms/AzureEssentialsPartners)